

The Canon Issue — Artometrics

The Canon Issue

How lists become memory — Gutenberg, Pantheon, Rolling Stone, and literary fame.

Which Subjects Dominate the Project Gutenberg Shelf?

Project Gutenberg is the internet's oldest large-scale experiment in free literature. Its subject catalog — Library of Congress headings and related labels — is a map of which public-domain books volunteers could digitize and classify at scale.

A TidyTuesday working extract of **100,000** records shows **lcsh** as the most common subject type, with tens of thousands of rows in that bucket. Subject codes such as **PS** recur thousands of times. The catalog is not every book ever written; it is the reusable shelf that copyright expiry, volunteer labor, and classification systems made easy to share.

Reading Gutenberg as a popularity contest misses the point. Reading it as a concentration map of reusable canon is where the data help.

Fast facts

The numbers that set the scale for this report:

100,000Records in the working dataset

lcshMost common Subject type

Data and method

The source is the TidyTuesday Project Gutenberg release from the R for Data Science community. The working file contains 100,000 rows after assembly — subject types, subject labels, and related catalog metadata.

Because many fields are categorical, the analysis leans on counts and repetition rather than a single quality score. Charts export as Plotly JSON with PNG fallbacks. Subject headings are librarian infrastructure, not reader reviews.

Public-domain books cluster by subject type

Public-domain books cluster by subject type

lcsh dominates with on the order of **69,027** records in the highlighted cut, far ahead of thinner subject-type buckets. The main classification family carries the story; this field does not split into many equal long-tail types.

That concentration means most navigational claims about what Gutenberg contains are really claims about how LCSH-style headings organize the digitized shelf.

A small set of subjects anchors the catalog

A small set of subjects anchors the catalog

PS appears about **4,684** times — among the most recurring subject names in the file. The top dozen account for a visible share of all 100,000 rows even though most subject entities appear only once.

American literature codes, fiction labels, and related headings form a reusable core. They are the catalog's gravitational center for teachers, scrapers, and adaptation hunters alike.

Subject families show the catalog center of gravity

Subject families show the catalog's center of gravity

lcsb is again the largest bucket on the category chart. Subject families show where editorial attention should focus first if the goal is to understand the shelf's center of gravity rather than its exotic edges.

The edges still matter for discovery. They do not define the statistical middle of a 100,000-row extract.

Repeated subjects reveal the reusable canon

Repeated subjects reveal the reusable canon

Most subject entities appear only once; a small head revisits repeatedly. That power-law shape is typical of catalog tables: a reusable canon of headings, and a long inventory of singleton classifications.

Repeated subjects are the ones most likely to support classroom packs, themed collections, and machine-learning corpora. Frequency is a reuse forecast as much as a shelf description.

Subject labels become the map of the shelf

Subject labels become the map of the shelf

PS and related labels become the map of the shelf when numeric scores are sparse. Frequency leaders reveal franchise depth in literature the way studio logos reveal franchise depth in film.

The practical claim for cultural analytics is simple: if you can only afford to study a slice of Gutenberg, the repeated subject head is where coverage of the reusable public-domain canon begins.

What this file cannot tell you

Community-cleaned TidyTuesday snapshots are not live APIs. Missing values, spelling variants, and sampling limits apply. Subject headings are not sales, downloads, or reading-time proof.

Findings describe structural signals about Project Gutenberg subject metadata — not a complete history of world literature, and not a ranking of artistic merit.

What to take away

Gutenberg's subject catalog is concentrated: lcsh dominates subject type, a short head of codes such as PS recurs thousands of times, and most labels appear once.

The citable lesson is about reusable canon. Public-domain literature becomes infrastructure when classification and digitization make the same subjects easy to find again and again.

Sources

Mapping the Public-Domain Canon Through Project Gutenberg

Project Gutenberg is not only a library. It is a map of what digitized public-domain culture looks like when volunteer labor, copyright timers, and language markets decide what is easy to share.

Public-facing summaries place the collection on the order of **75K+** ebooks, with catalog feeds updating on **weekly** CSV and **daily** RDF cadences in Gutenberg tooling docs. The first chart compares **8** languages; later panels use **10** authors as anchors for availability and memory.

English gravity, nineteenth-century density, and adaptation-ready subjects are the three forces that turn a free shelf into a cultural operating system for schools, remixers, and machines.

Fast facts

The numbers that set the scale for this report:

75K+ Approximate scale of Project Gutenberg ebooks in public-facing summaries

Weekly CSV catalog update cadence noted in Gutenberg tooling docs

Daily RDF catalog update cadence noted by Gutenberg

8 Languages compared in the first chart

10 Authors used as anchors

Data and method

Project Gutenberg provides machine-readable metadata in RDF/XML, MARC, and CSV formats. Official pages recommend using the metadata feeds rather than scraping HTML. A scaled production pass should ingest the weekly CSV or RDF catalog, normalize subjects, languages, and authors, and join to adaptation, syllabus, and Wikidata signals.

The charts here use editorial indices to define that comparative structure before a full catalog join. They are meant to make language gravity, era mix, author memory, and remix power legible — not to replace the official feeds.

Language gravity

English dominates the accessible public-domain shelf

English dominates the accessible public-domain shelf in the comparison of eight languages — on the order of **72** on the index scale used in the chart, with French and German far behind and Finnish, Dutch, Portuguese, Italian, and Spanish in a thinner band.

Language gravity is not a judgment about literary quality. It is a statement about which texts are easiest to search, remix, teach, and feed into downstream culture when copyright and digitization align.

The era machine

The 19th century becomes the public-domain literary core

Copyright law turns time into a cultural filter. Fiction and literature indices rise sharply through the nineteenth century; nonfiction and reference follow a related but not identical path. Pre-1800 and post-1950 bins look thinner on the reusable shelf.

That makes Gutenberg a meso dataset: not every book ever written, but the books most ready to be reactivated because they are old enough, popular enough, and digitized enough.

Author memory

Digital availability and cultural memory reinforce each other

Authors become infrastructure when their works are everywhere: classrooms, editions, audiobooks, adaptations, quote databases, and training corpora. Digital availability and cultural memory reinforce each other rather than acting as independent clocks.

That feedback loop is why public domain matters to AI culture too. What is easy to load is what models and people alike can easily learn from.

Adaptation power

Adventure, childhood, and gothic subjects adapt especially well

Adventure, manners, gothic, science, war, travel, religion, and children's subjects do not adapt equally. Adventure, childhood, and gothic titles move especially cleanly across film, brand memory, and genre grammar.

The public-domain market is therefore not dead literature. It is reusable cultural material with uneven conversion rates into new media.

Canon and remix

Some books become curriculum while others become remix engines

Some books become curriculum while others become remix engines. *Pride and Prejudice* is classroom canon and adaptation fuel. Sherlock Holmes, Dracula, and Frankenstein behave more like cultural APIs: infinitely forkable characters and premises.

The strongest public-domain works are not only read. They are reused — assigned, quoted, adapted, and recombined until availability becomes a kind of fame.

What to take away

The public-domain canon is not neutral. It is shaped by language, copyright age, volunteer digitization, educational reuse, and adaptation economics.

Gutenberg is a bridge dataset: it connects literature, AI training culture, education, film adaptation, and historical memory through the same shelf of reusable texts.

Sources

Who Gets Remembered in the Global Pantheon of Fame?

Collective memory is not a neutral archive. It is a ranked attention system: some lives travel across languages and centuries, while others remain local even when their achievements were large.

Pantheon-style biography indexes — on the order of **85K+** lives in public-facing summaries, with a collection logic that requires presence in at least **15** Wikipedia language editions — make that ranking measurable. The 2016 *Scientific Data* paper for Pantheon 1.0 is the scholarly anchor. The charts below treat three main memory engines — power, art, and science — across **10** memory-capital cities.

If humanity remembers some kinds of people more than others, that pattern is itself a cultural object. The question is which occupations and cities compound fame, and which leave gaps.

Fast facts

The numbers that set the scale for this report:

85K+Biographies indexed by Pantheon in public-facing summaries
15Minimum Wikipedia language editions used in Pantheon collection logic
2016Scientific Data paper for Pantheon 1.0
3Main memory engines: power, art, science
10Memory-capital cities compared

Data and method

Pantheon aggregates biographical notability using multilingual Wikipedia presence and related fame metrics. The dataset is powerful because cross-language edition counts act as a rough passport stamp for global memory — a person who exists in many Wikipedias has traveled farther through culture than a person who exists in one.

The charts here use editorial indices to make domain, era, city, and occupation comparisons legible. They are interpretive maps grounded in Pantheon's public logic, not a claim that Wikipedia is the only archive that matters.

Memory domains

Power, art, and science travel farther across collective memory

Politics, arts, science, sports, business, religion, exploration, and activism do not travel equally through collective memory. Power, art, and science repeatedly show up as the main engines that cross languages and centuries.

Sports and business can be enormously famous in the present and still underperform on long multilingual durability. The domain chart separates contemporary visibility from archival carry.

Era shift

The remembered person changes from ruler to cultural or scientific producer

Across eras — ancient, medieval, early modern, 1800s, 1900s, 2000s — the remembered person changes. Rulers dominate earlier memory shares; cultural and scientific producers rise as print, universities, and later mass media create new fame factories.

The shift is not a morality play about progress. It is a change in which institutions mint durable names: courts and churches first, then academies and culture industries.

Memory capitals

Cities become archives when institutions compound across centuries

Paris, London, Rome, New York, Vienna, Berlin, Athens, and Moscow appear among the memory-capital comparisons because institutions compound there: universities, courts, publishers, studios, academies. Cities become archives when those institutions last long enough to keep minting and storing notable lives.

A memory capital is not identical to a current GDP capital. Some cities punch above their present economic weight because their historical institutions still feed the multilingual record.

Two clocks of fame

Some occupations have durable translation while others spike through present attention

Some occupations show durable translation across languages; others spike through present attention without the same archival carry. Literature, film, music, science, sports, and politics sit on different clocks — one measured in centuries of citation, one measured in contemporary noise.

The two-clock chart is the practical tool for editors and educators: it warns against treating a viral profession as automatically canonical, and against treating a canonical profession as automatically popular.

Archive gap

The archive remembers through historical inequality

The archive remembers through historical inequality. Domains and regions with denser institutions, denser literacy, and denser Wikipedia coverage look more important because they are more recorded.

That gap is not a reason to discard Pantheon. It is the reason to read it as a map of remembered inequality — a cultural object that reveals what global memory systems have been able to store.

What to take away

Collective memory, measured through Pantheon-style multilingual biography indexes, concentrates on power, art, and science, compounds in a short list of memory-capital cities, and runs on two clocks of fame.

The citable claim is modest: what humanity remembers at global scale is structured — by occupation, era, and city — and that structure is itself evidence about cultural power.

Sources

Which Albums Sit at the Peak of Rolling Stone's Canon?

Rolling Stone's greatest-albums lists are canon factories. They convert critical memory into ranked cultural infrastructure — and the 2003 ranking spine in this TidyTuesday extract is one of the most cited versions of that factory.

A working file of **691** records puts median 2003 rank at **250** and the highest observed rank number at **500**. **Touch** leads the album ranking by the 2003 rank field in the fact box, the year span of album release years runs **1955–2019**, and **Studio** is the most common type label.

Remember the rank inversion: lower numbers are greater prestige on a traditional list, while some chart cuts sort the numeric field as a magnitude. The prose below follows the file's stated leaders and medians without inventing a second scoring system.

Fast facts

The numbers that set the scale for this report:

691Records in the working dataset
250Median Rank 2003
500Highest observed Rank 2003
TouchTop Album by Rank 2003
1955–2019Year span covered in the file
StudioMost common Type

Data and method

The source is the TidyTuesday Rolling Stone albums release from the R for Data Science community. The working file contains 691 rows after cleaning — album titles, artists, release years, types, genre-family labels, and 2003 rank fields among them.

Medians summarize a list that is already a curated elite rather than a random sample of recorded music. Charts export as Plotly JSON with PNG fallbacks. A magazine canon is an editorial object; treating it as market share will misread the politics of taste.

How ranks sit across release years

Median Rank 2003 Over Time

Median 2003 rank across release years shows which eras the list's taste machine favored when the ranking was locked. Mid-century and classic-rock dense decades often sit differently from later indie and hip-hop entries depending on how the canon was composed.

Because the rank number itself comes from a 2003 list, later release years in the file are a reminder that the table can carry albums that post-date or sit awkwardly against that freeze-frame — coverage quirks included.

Who sits at the top of the numeric rank field

Touch leads at 500 — 494 marks the median among the top dozen

Touch leads at **500** on the highlighted cut, with **494** as the median among the top dozen. In this chart's ordering, the upper numeric band is a cluster of albums near the far end of the 2003 list's numbering.

Nearby titles in related leader cuts — *Guitar Town*, *Entertainment!*, *All the Young Dudes*, *Vitalogy*, *She's So Unusual* — show how the list's lower prestige tier is still a hall of fame by ordinary-market standards. Being near 500 on a *Rolling Stone 500* is not obscurity.

Ranks by album type

Rank 2003 by Type

Box plots by type — Studio foremost, with other type labels where present — show whether rank consensus is shared or contested across release formats. Studio albums dominate headcount because rock-critical canons historically privileged them.

Type concentration is an editorial fingerprint. A canon that underweights compilations, live albums, or other forms is making an aesthetic argument, not merely counting what exists.

Who beats the median — and who trails

Rank 2003 vs median by Type

The gap chart ranks types above or below the median 2003 rank. Differences encode which formats the list treated as central versus peripheral when the ranking was composed.

Because the whole file is already an elite list, trailing the median still means belonging to a magazine's chosen 500 — a narrow world compared with all recorded music.

Rank and artist age

Rank 2003 vs Ave age at top 500

Plotting 2003 rank against average age at the top-500 marker shows clusters that averages erase. Some high-prestige albums come from young breakthrough artists; others from later-career peaks. The scatter refuses a single myth about youth and canonization.

Age is not destiny on a critical list, but it is a recurring covariate in rock historiography. Seeing the cloud is better than assuming every classic was made at 23.

What this file cannot tell you

Community-cleaned TidyTuesday snapshots are not live APIs. Missing values, spelling variants, and list-version mismatches apply. A 2003 rank is not a 2020s streaming share, and Rolling Stone's taste is not the world's taste.

Findings describe structural signals about one magazine canon's metadata — not a universal ranking of recorded music, and not a substitute for listening.

What to take away

Rolling Stone's album canon in this file is a curated elite of 691 rows with a median 2003 rank of 250, studio albums as the default type, and a far-end numeric cluster near 500 that is still culturally famous.

The citable caution is genre and generation politics: canons freeze taste at a moment, then keep teaching that freeze as if it were nature. The charts make the freeze measurable.

Sources

How Word Count Is Distributed Across the Holmes Canon

Conan Doyle's Sherlock Holmes corpus arrives here as a tokenized text table: 65,958 records and a median word count of 12.0. That median is the calibration point for everything that follows — not novel length in the bookstore sense, but the word-count field as stored in this working extract.

The Adventure of the Beryl Coronet tops the fact-box ranking by word count, while the charted story leaders highlight The Yellow Face at 13.0 against a top-dozen median of 12.0. The canon is famous for plot; this file is interested in textual mass and how that mass concentrates across books.

Fast facts

The numbers that set the scale for this report:

65,958Records in the working dataset

12.0Median Word count

18.0Highest observed Word count

The Adventure of the Beryl CTop Book by Word count

Data and method

The source is the TidyTuesday release from 2025-11-18 (R for Data Science community). The working file contains 65,958 rows and 4 columns after merging available tables in the week folder. Book titles are the primary categorical axis; word count is the primary numeric metric.

Medians are used because text-length fields can skew. The distribution chart also reports a mean of 10.9 against the median of 12.0, a relatively symmetric shape for this extract. Index-style fields stay out of metric selection.

Story-level medians sit in a narrow band

Word count by Book

In the book breakdown, *The Yellow Face* leads at 13.0 while *A Case of Identity* anchors the low end of the plotted group at 12.0. The spread among visible leaders is small — a reminder that this word-count field clusters tightly across many titles.

Grouping by book exposes how the metric varies across the catalog's major entities without requiring plot summary. The question is length structure, not detective method.

Even the leaders barely clear the file median

The Yellow Face leads at 13.0 — 12.0 marks the median among the top dozen

The Yellow Face leads at 13.0, and 12.0 marks the median among the top dozen — identical to the file-wide median of 12.0. In this extract, leading does not mean a dramatic leap above the typical observation; it means a modest edge inside a compressed distribution.

That compression is itself a finding. Holmes stories in this tokenization do not fan out into wildly different word-count regimes at the center of the leaderboard.

Median 12.0, mean 10.9 — a mostly symmetric shape

Word count Distribution

The full-sample distribution shows median 12.0 versus mean 10.9 — relatively symmetric rather than a heavy right skew. The top decile begins at 14.0; that upper slice is where the longer textual cases in this metric live.

Symmetry matters for citation. A median-centered story about Holmes text length is more faithful here than an average pulled by a few extreme tokens.

Five books account for a third of aggregate word count

Cumulative Word count

The top five book entries account for 34% of the aggregate word count. Concentration is real but not absolute: a meaningful head exists, and a long tail of other titles still holds most of the remaining mass.

Pareto steepness here means editorial attention can start with a small set of books without pretending the rest of the canon is empty.

A second concentration cut confirms the same head

Cumulative Word count

A second cumulative chart repeats the concentration view: again, the top five book entries account for 34% of aggregate word count. Parallel Pareto cuts are useful as a robustness check when the pipeline exports more than one concentration figure.

The takeaway does not change. A compact set of titles carries a disproportionate share of the summed word-count field in this 65,958-row extract.

What this file cannot tell you

Community-cleaned TidyTuesday snapshots are not live literary corpora APIs. Tokenization choices, spelling variants in titles, and week-of-export coverage limits apply. Word count here is the field in the file — not an independent scholarly page count.

Findings describe this Holmes extract. They are structural signals about textual mass across books, not claims about literary quality or cultural importance.

What to take away

Holmes in this file is a tightly centered length distribution: median word count 12.0, mean 10.9, top decile from 14.0, and leaders such as *The Yellow Face* only a step above the center.

Concentration still matters — the top five books hold 34% of aggregate word count — but the drama is less runaway novels than a canon whose tokenized mass piles up without extreme median drift.

Sources

Data Science Learning Community. (2025). *TidyTuesday: Sherlock Holmes*. <https://raw.githubusercontent.com/rfordatascience/tidyuesday/main/data/2025/2025-11-18/holmes.csv>

Editor's note

Artometrics data report from the TidyTuesday research pipeline. Charts and aggregates are reproducible from the embedded exhibits and public source files.

[View TidyTuesday source on GitHub](#)

Which Christmas Novels Anchor the Holiday Shelf?

Christmas novels are a seasonal shelf with a long memory: Dickensian ghosts, children's mysteries, parlor sentimentalism. The TidyTuesday christmas-novels extract used here is small — 35 records — with a median author birthdate of 1,859 and a latest birthdate of 1,891. Charles Dickens is the most common author label in the file.

Because the sample is tiny, every chart is a map of a curated canon rather than a complete market. Birthdate becomes a proxy for which literary generations stock the holiday shelf.

Fast facts

The numbers that set the scale for this report:

35Records in the working dataset
1,859Median Birthdate
1,891Highest observed Birthdate
The Camp Fire Girls Solve a Top Title by Birthdate
1970–1970Year span covered in the file
Dickens, CharlesMost common Author

Data and method

The source is the TidyTuesday release from 2025-12-30 (christmas_novels.csv). After cleaning, 35 rows remain.

Author birthdate is the primary numeric metric in this chart stack; deathdate appears in the joint plot. Charts are Plotly JSON with PNG fallbacks.

The Shelf Centers on Mid-Nineteenth-Century Births

Median Birthdate Over Time

The trend summary in this build sits at a median birthdate of **1859** — Victorian and immediately post-Victorian authors anchoring the holiday novel list.

That center of gravity explains the tone of the canon: industrial-era childhood, Christian seasonal ritual, and serialized storytelling habits.

Later-Born Titles Mark the Edge of the Canon

The Camp Fire Girls Solve a Mystery; Or, The Christmas Adventure at Carver House leads at 1,891 — 1,872 marks the median among the...

The Camp Fire Girls Solve a Mystery; Or, The Christmas Adventure at Carver House leads at birthdate **1891**. Other later births include Uncle Noah's Christmas Inspiration (**1884**) and Christmas Outside of Eden (**1883**). The median among the top dozen by birthdate is about **1872**.

Even the “late” edge of this shelf is still nineteenth-century. Modern holiday romances are mostly outside this particular extract.

Birthdates Cluster Around the 1860s

Median 1,859 vs mean 1,849 — the shape is relatively symmetric

The distribution piles several authors near the early 1860s bins, with smaller counts in earlier nineteenth-century bands. Mean birthdate (~**1849** in the chart annotation's neighborhood) sits close enough to the median to read as relatively symmetric for such a small n.

With only 35 rows, bin heights are fragile — but the Victorian concentration is robust enough to cite.

Dickens and Thackeray Sit Earlier Than the Median

Birthdate vs median by Author

Author gaps to the median birthdate show Rupert Hughes furthest above (+13), while Dickens (−47) and Thackeray (−48) sit well earlier. Louisa May Alcott is also clearly before the median (−27).

The holiday shelf's most famous names are often older than the shelf's statistical middle — canon gravity pulling toward earlier births.

Birth and Death Dates Trace Finite Author Lives

Birthdate vs Deathdate

Birthdate versus deathdate for 32 plotted authors follows the expected upward diagonal: later-born authors die later, with points stretching from early-nineteenth-century pairs into mid-twentieth-century deaths.

The scatter is a mortality timeline for a seasonal literature, not a sales chart — useful for placing which generations could still be writing when Christmas publishing industrialized.

What this file cannot tell you

n=35 is a teaching canon, not a market census. Missing modern titles and non-English traditions are structural absences. Birthdate as a ranking metric is biographical, not a measure of holiday popularity.

Author name authority control may still split or merge identities in a small file.

What to take away

This Christmas-novel shelf is a Victorian-centered canon: median birth-date 1859, latest 1891, with Dickens as the recurring name.

Cite the small n alongside any ranking. The value is literary geography — which generations stock the holiday imagination — not a bestseller index.

Sources

Data Science Learning Community. (2025). *TidyTuesday: Christmas Novels*. https://raw.githubusercontent.com/rfordatascience/tidyuesday/main/data/2025/2025-12-30/christmas_novels.csv

Did Longer Medium Posts Earn More Applause?

Medium promised a writing network where attention could attach to length, craft, and niche expertise. The metadata trail — reading time, claps, publication tags — lets that promise be tested without reading every essay.

A TidyTuesday extract of **78,388** article rows from **2017–2018** puts median reading time at **4.00** minutes. The highest observed reading time hits **100** minutes — a month-long chess-mastery quest that sits at the extreme of the length distribution. **Towards Data Science** is the most common publication label in the file.

The open question is whether longer posts earned more applause. Reading time and claps are related enough to chart, and messy enough that a single correlation will not settle the culture of the platform.

Fast facts

The numbers that set the scale for this report:

78,388Records in the working dataset
4.00Median Reading time

100Highest observed Reading time
My month-long quest to becomeTop Title by Reading time
2017–2018Year span covered in the file
Towards Data ScienceMost common Publication

Data and method

The source is the TidyTuesday release from 2018-12-04 (R for Data Science community). The working file contains 78,388 rows and 22 columns after assembly — titles, publications, reading time, claps, and related metadata.

Medians stabilize a distribution with a long right tail of mega-posts. Charts export as Plotly JSON with PNG fallbacks. Reading time is a platform estimate, not a stopwatch on every reader.

How reading time moved

Median Reading time Over Time

Across the short 2017–2018 window, median reading time stays near **4.00** minutes from opening period to close. Stability at the median does not mean the tails were quiet — only that the typical post length did not reprice during the snapshot.

A flat median is still informative. It suggests that whatever clap dynamics existed in this period, they were not driven by a wholesale shift toward longer or shorter default essays.

Who sits at the top of length

My month-long quest to become a chess master from scratch leads at 100 — 68.0 marks the median among the top dozen

My month-long quest to become a chess master from scratch leads at **100** minutes of estimated reading time. The median among the top dozen is **68.0** — still more than fifteen times the file-wide median of 4 minutes.

Deep-learning lesson writeups, AI alignment podcasts, and longform data essays populate the same extreme band. The length frontier of Medium in this file is a specialist pedagogy genre as much as a literary one.

How publications spread length

Reading time by Publication

Publication box plots show whether reading-time consensus is shared or contested across houses. Some publications cluster tightly around short explainers; others tolerate or encourage long technical serials.

Towards Data Science's frequency as the most common publication does not automatically make it the longest. Volume of posts and length of posts are different editorial strategies that can coexist under one brand.

Who beats the median — and who trails

Reading time vs median by Publication

Towards Data Science sits **2.00** minutes above the median on the gap chart; **Data Driven Investor** trails by **0.00** in the highlighted comparison — effectively on the median line in that cut.

Publication gaps of a couple of minutes sound small until you remember the median is only four. A two-minute lift is a 50% longer typical article in that house's mix.

Reading time and claps

Reading time vs Claps

Plotting reading time against claps shows clusters that averages erase. Many short posts earn modest applause; some long posts earn a lot; some long posts earn almost none. Length is not a reliable applause machine.

If there is a relationship, it is noisy and genre-dependent. Tutorial serials and viral short takes can both succeed. The scatter's job is to kill the slogan that longer always wins without replacing it with the opposite slogan.

What this file cannot tell you

Community-cleaned TidyTuesday snapshots are not live APIs. Missing values, spelling variants, and 2017–2018 coverage limits apply. Claps are a platform-specific applause metric, not revenue or unique readers.

Findings describe structural signals about Medium article metadata in the release window — not a complete theory of online writing incentives after paywalls, algorithms, and audience shifts.

What to take away

Most Medium posts in this file are short: median reading time of four minutes. A thin upper tier stretches to hour-scale essays, especially in technical and quest narratives.

The citable answer to the title question is cautious: longer posts sometimes earn more claps, but the scatter refuses a clean law. Publication norms and genre explain as much as length alone.

Sources