

The Culinary Issue — Artometrics

The Culinary Issue

Coffee, beer, wine, ramen, and the prestige economics of taste.

Coffee's Place in Global Trade and Daily Life

Coffee is the world's second-largest traded commodity by value, behind only crude oil. More than 125 million people depend on it for their livelihoods. Around 2.25 billion cups are consumed every day. And yet the country that produces the highest-quality coffee on earth — by a significant margin, according to the Coffee Quality Institute's own grading data — has exactly zero Starbucks locations. That country is Ethiopia.

What follows uses CQI cupping scores and country-level samples to ask a simple question: where does quality concentrate, and how little that map resembles the retail geography of coffee brands.

Research question

Does the geography of high-scoring coffee match the geography of retail coffee power? This report asks whether Coffee Quality Institute cupping scores and Starbucks store counts describe the same global map or two different positions in the commodity chain.

The observational question is where Arabica quality concentrates once country medians and sensory sub-metrics are plotted, and whether that origin map aligns with the markets where a dominant branded re-

tailer operates stores. The answer matters because origin value, sensory quality, and retail capture are often discussed as if they naturally reinforce one another.

Fast facts

Keep these markers in view as the story unfolds.

1,339 Coffee samples rated by CQI Q Graders across Arabica and Robusta species

85.3 Ethiopia's median CQI cup score — highest of any country with 20+ samples

0 Starbucks locations in Ethiopia, Kenya, and Uganda — the top three scoring nations

16 Countries with 20+ CQI-rated Arabica samples, spanning four continents

7 Sensory sub-metrics scored per sample — aroma, flavor, aftertaste, acidity, body, balance, cupper points

25,600 Starbucks locations worldwide across 73 countries, as of the 2018 dataset

Data and method

The Coffee Quality Institute is a nonprofit that trains and certifies Q Graders — licensed sensory evaluators who score coffee using a standardized 100-point protocol. A sample must score 80 or above to qualify as specialty coffee. The CQI dataset made available through TidyTuesday (2020-07-07) contains **1,339** samples spanning Arabica and Robusta species, with individual scores across ten sub-metrics plus a total cup points aggregate. This analysis focuses exclusively on Arabica. After removing one data entry error (a near-zero-score row), the working Arabica dataset contains 1,312 samples from 36 countries.

The Starbucks locations dataset (TidyTuesday 2018-05-07) documents **25,600** individual store locations across 73 countries, with brand, ownership type, city, and geographic coordinates. For the quality-versus-retail chart, Starbucks counts were aggregated at the country level and joined to CQI quality scores using ISO-2 country codes.

The sub-metric fingerprint covers seven of the ten CQI sub-metrics: aroma, flavor, aftertaste, acidity, body, balance, and cupper points. Uniformity, clean cup, and sweetness were excluded because they function

more as defect-penalty fields than sensory attributes. The top eight countries by sample volume were selected for the heatmap to ensure statistically stable medians.

Where Great Coffee Comes From

Country Ridgeline

Ethiopia doesn't just lead — it occupies a different part of the distribution entirely. While the median cup score across all 16 qualifying countries sits at 82.5, Ethiopia's distribution peaks near 85 and has a long right tail that no other origin comes close to matching. Kenya and Uganda sit in second and third, but their distributions overlap substantially with the rest of the field. Ethiopia's does not.

What makes this finding significant is what it implies about terroir and altitude. Ethiopia is the genetic origin of Arabica coffee — the Kaffa region in southwestern Ethiopia is where the species was first documented growing wild. The country's range of elevation, soil diversity, and processing tradition produces a flavor complexity that Q Graders consistently reward with above-median scores across every sub-metric.

The bottom of the chart is equally instructive. Nicaragua, Mexico, Honduras, and Taiwan all cluster between 81 and 83 — technically specialty grade, but right at the floor. Brazil, the world's largest coffee producer by volume, lands firmly in the lower half — a reminder that market dominance and cup quality are not the same thing.

The Quality-Retail Disconnect

Quality Vs Retail

There is no positive relationship between a country's quality score and Starbucks' retail presence there. The three highest-scoring origins — Ethiopia, Kenya, and Uganda — all sit at zero Starbucks locations. Meanwhile, Hawaii (USA) lands near the top of the y-axis not because of its quality rank alone but because its ISO code maps to the US, which hosts thousands of Starbucks locations.

This is structural. Starbucks builds stores in markets where it can sell coffee, not only in markets where it sources coffee. Ethiopia's retail consumer market is too small and too low-income to support the Starbucks price point at scale. The company sources Ethiopian beans, but that sourcing relationship does not translate into a domestic retail presence.

Colombia reinforces the pattern: a recognized origin with specialty scores that still shows little Starbucks store footprint relative to the fame of its beans. Economic value from growing specialty coffee does not automatically accumulate where the coffee is grown.

The Sub-Metric Fingerprint

Submetric Heatmap

Ethiopia doesn't just score higher overall — it scores higher on every single sub-metric among the eight-country heatmap field. Its median acidity, cupper points, and flavor lead the set. Kenya comes close on aroma and flavor but trails on most other dimensions.

What the heatmap also reveals is how flat the remaining countries are. Brazil sits in a narrow mid band across attributes — consistent and undifferentiated. Honduras and Mexico are nearly indistinguishable on several scores. When buyers talk about commodity coffee, this is the sensory data behind the phrase.

Guatemala and Colombia show elevated acidity and balance relative to that flat middle — profiles that translate easily into consumer language about brightness and approachability.

What this file cannot tell you

The CQI dataset reflects what Q Graders evaluate — not a random sample of all coffee produced. Farms that can afford evaluation tend to be larger or already tied into specialty chains, so subsistence-grade output is underrepresented.

The Starbucks dataset is from 2018. Store counts should not be treated as current, even if the directional finding about missing retail presence in top African origins has remained directionally stable. Starbucks is a proxy for one dominant chain, not the full café and grocery landscape.

Country-level analyses were filtered to origins with 20 or more CQI samples, cutting the pool from 36 to 16 countries. Prestigious origins with fewer submissions — including some ultra-premium varieties — are absent from the heatmap.

What to take away

Quality and market power are not the same thing, and they do not flow in the same direction. Ethiopia produces the highest-rated coffee in this CQI slice by a significant margin and shows zero Starbucks locations in the joined retail file.

Brazil dominates volume while scoring in the lower half of this quality ranking. The countries that grow the best-scoring coffee are not automatically the countries that capture the most retail value from coffee's global brand.

Sources

The Coffee Quality Institute data should be interpreted as a specialty-coffee evaluation archive, not a random sample of all coffee production. CQI Q Graders score submitted samples under standardized protocols, which means the file is strongest for comparing evaluated specialty lots and weaker for estimating national production averages. The Starbucks file adds a retail-chain geography from 2018, useful as a branded-store proxy but not as a full measure of cafe, grocery, roaster, or export-market value.

The broader citation frame comes from coffee-development and value-chain research by Daviron, Ponte, Samper, Quiñones-Ruiz, and the International Coffee Organization. Those sources matter because the charts show a classic commodity-chain split: Ethiopia, Kenya, Uganda, Colombia, Brazil, and other origins can produce sensory value, while roasting, branding, retail real estate, and consumer pricing often capture value elsewhere.

TidyTuesday (2020-07-07). *Coffee Ratings*. R for Data Science Community. github.com/rfordatascience/tidytuesday/.../2020-07-07

TidyTuesday (2018-05-07). *Coffee Chains*. R for Data Science Community. github.com/rfordatascience/tidytuesday/.../2018-05-07

Coffee Quality Institute. (2020). *Q Grader Program and Cupping Protocols*. Retrieved from coffeeinstitute.org

Samper, L. F., & Quiñones-Ruiz, X. F. (2017). Towards a balanced sustainability vision for the coffee industry. *Resources*, 6(2), 17. doi.org/10.3390/resources6020017

Daviron, B., & Ponte, S. (2005). *The Coffee Paradox: Global Markets, Commodity Trade and the Elusive Promise of Development*. Zed Books.

Wintgens, J. N. (Ed.). (2009). *Coffee: Growing, Processing, Sustainable Production* (2nd ed.). Wiley-VCH.

Editor's note

This analysis was researched, written, designed, and produced in active collaboration with Claude AI (Anthropic). The data pipeline, statistical analysis, chart design, written analysis, narrative structure, and visual styling were all developed through a directed partnership between human editorial judgment and AI execution.

Artometrics was built on the premise that rigorous analysis and honest process are not in conflict. The research questions, editorial instincts, interpretive framing, and brand vision are ours. The execution — every line of R code, every paragraph of analysis, every design decision — was a collaboration. We document this not as a disclaimer but as a description of how we actually work, and as a position: we believe this is what serious data journalism looks like when the tools available are used honestly and at full capacity.

— Artometrics Editorial

[View source on GitHub](#)

How Bitter and Strong Is American Craft Beer?

American craft beer advertises itself with bitterness and strength — IBU and ABV on the tap list. The TidyTuesday beers extract used here holds **2,410** records with a median ABV of **0.06** (6%) and a high of **0.13** (13%) for Lee Hill Series Vol. 5 - Belgian Style Quadrupel Ale.

Most craft beers in this file are not barrel-aged monsters. They cluster near session-to-standard strength, with a right tail of barleywines, quads, and imperial experiments.

Fast facts

The numbers that set the scale for this report:

2,410Records in the working dataset

0.06Median Abv

0.13Highest observed Abv

Lee Hill Series Vol. 5 - BelTop Name by Abv

Data and method

The source is the TidyTuesday release from 2018-07-10 (week15_beers.xlsx). After cleaning, 2,410 rows remain.

ABV is the primary ranked metric; IBU joins in the scatter. Charts are Plotly JSON with PNG fallbacks.

High-ABV Outliers Are Named Experiments

Abv by Name

Lee Hill Series Vol. 5 leads at **0.128** ABV, with London Balling at **0.125** and Csar at **0.120**. Other leaders — Lee Hill Vol. 4, 4Beans, Johan the Barleywine, Wizard Burial Ground, Hi-Res — sit near **0.10**.

These titles advertise intensity. They are the marketing edge of craft, not the median pint.

The Top Dozen Still Doubles Everyday Strength

Lee Hill Series Vol

The leaders remain capped by the 0.128 quad, with the rest of the top dozen occupying the 0.10–0.12 band — roughly double the 0.06 median.

Cite 0.06 as what “typical craft ABV” means in this extract, and 0.10+ as the specialty high-gravity club.

Most Beers Pile Up Near 5–6% ABV

Median 0.06 vs mean 0.06 — the shape is right-skewed

The distribution's tallest bins sit around roughly **0.05–0.06** ABV (hundreds of beers each), with counts falling as ABV climbs toward 0.09–0.10 and almost vanishing above 0.11.

Right skew is real but thin: the famous 12–13% outliers are rare relative to the pale-ale middle.

High-ABV Names Concentrate in a Short List

The top 5 name entries account for 37% of the aggregate abv

The Pareto view of leading high-ABV names shows the first five entries holding about **37%** of the plotted aggregate strength share, rising toward the full total by the fifteenth name.

Extreme ABV is a boutique phenomenon inside the craft catalog — concentrated among a few loudly strong beers.

Bitterness and Strength Rise Together — Imperfectly

Abv vs Ibu

ABV versus IBU across roughly 1,400 plotted points shows many high-IBU IPAs in the 0.06–0.09 ABV band, with some of the bitterest beers exceeding 120 IBU. High ABV does not guarantee high IBU; stouts and quads can be strong without chasing hop extremes.

The joint distribution is the craft-beer design space: hop-forward versus malt-forward paths to intensity.

What this file cannot tell you

ABV and IBU fields can be missing or rounded. Style coverage depends on which breweries entered the scrape. “Craft” definitions and the 2018 vintage miss later hard-seltzer and nonalcoholic shifts.

Name-level leaders may be one-off specialty releases, not flagship year-round beers.

What to take away

In 2,410 U.S. craft beers, median ABV is 0.06 while the ceiling reaches 0.128. Most of the catalog lives near everyday strength; the high-gravity tail is short and famous.

Cite the ABV–IBU scatter when explaining style strategy: bitterness and alcohol are correlated tools, not a single dial.

Sources

Data Science Learning Community. (2018). *TidyTuesday: Craft Beer USA*. https://raw.githubusercontent.com/rfordatascience/tidytuesday/main/data/2018/2018-07-10/week15_beers.xlsx

Does Wine Price Predict Critic Points?

A hundred thousand wine ratings are enough to test a dinner-table theory: does price predict critic points? This file’s median score is 88.0; the observed high is 100. Charles Smith 2006 Royal City appears in the fact-box title ranking; the United States is the most common country.

The charts move from title leaders to country spreads to the points-versus-price scatter. The calibration point is 88.0 — the mathematical center of a tasting culture that rarely uses the bottom of the scale.

Fast facts

The numbers that set the scale for this report:

100,000Records in the working dataset

88.0Median Points

100Highest observed Points

Charles Smith 2006 Royal CitTop Title by Points

USMost common Country

Data and method

The source is the TidyTuesday release from 2019-05-28 (R for Data Science community). The working file contains 100,000 rows and 13 columns after merging available tables in the week folder. Points is the primary metric; price appears in the scatter; title and country are categorical axes.

Medians are used because wine scores cluster tightly and prices skew. Index-style fields are excluded from metric selection.

Perfect scores crowd the title breakdown's top slice

Points by Title

In the title breakdown, Château Léoville Barton 2010 Saint-Julien leads at 100; Cardinale 2006 Cabernet Sauvignon (Napa Valley) anchors the low end of that plotted top group at 100 as well — a flat elite band at the scale's ceiling.

When the top slice is all 100s, the chart is telling you about saturation at the maximum, not about a wide ladder of near-misses.

The leaders chart is a wall of 100-point wines

Château Léoville Barton 2010 Saint-Julien leads at 100 — 100 marks the median among the top dozen

Château Léoville Barton 2010 Saint-Julien leads at 100, and 100 marks the median among the top dozen. The elite is compressed at the ceiling. Charles Smith 2006 Royal City remains a fact-box landmark for the title ranking rule used there.

A top dozen whose median is 100 is not a typical wine; it is the ceremonial head of a 100,000-row tasting archive.

Countries spread points around a high center

Points by Country

Category boxes by country show whether points consensus is shared or contested across origins. The United States is the most common country by frequency; the boxes ask how scores spread inside each origin.

Country structure matters because wine writing often treats place as destiny. The boxes test how much score variation place actually carries in this file.

Austria clears the median; Chile trails

Points vs median by Country

Austria sits 2.00 above the median; Chile trails by 2.00. Those signed gaps are modest in points terms — a reminder that this scoring culture lives in a narrow band around 88.0.

Small gaps can still be meaningful when the whole scale's action happens between the high 80s and the low 90s for most bottles.

Points and price form the real test of the dinner-table theory

Points vs Price

Plotting points against price shows clusters that averages erase: cheap high scorers, expensive middling bottles, and the familiar cloud where spending more often accompanies higher scores without guaranteeing them.

The scatter is the relationship the title promised. It does not invent a correlation coefficient beyond what the chart shows; it shows the joint distribution in the 100,000-row extract.

What this file cannot tell you

Community-cleaned TidyTuesday snapshots are not live Wine Enthusiast APIs. Missing prices, title variants, and coverage limits apply. Critic points reflect the source's scoring practice, which uses a compressed high end of the scale.

Findings describe this wine-ratings extract — structural signals about points, country, and price — not a universal law of taste.

What to take away

Wine scores in this file center on a median of 88.0, peak at 100 among title leaders, and show only modest country gaps around that high center.

Austria sits above the median and Chile below by two points each, while the points–price scatter is where the cost-versus-score question becomes visible without reducing 100,000 tastings to a slogan.

Sources

Data Science Learning Community. (2019). *TidyTuesday: Wine Ratings*. <https://raw.githubusercontent.com/rfordatascience/tidyuesday/main/data/2019/2019-05-28/winemag-data-130k-v2.csv>

Editor's note

Artometrics data report from the TidyTuesday research pipeline. Charts and aggregates are reproducible from the embedded exhibits and public source files.

[View TidyTuesday source on GitHub](#)

What Does a Slice of Pizza Really Cost?

Pizza prices look simple until you try to compare 10,000 menu rows. The Datafiniti pizza extract in TidyTuesday does exactly that: restaurant-level records with price-range fields that reveal how “a slice” and “a specialty pie” live in different economic brackets. The working file holds **10,000** rows.

The median price-range minimum is **0.00** — a signal that many listings encode missing or zeroed floors rather than free pizza. The highest observed price-range minimum is **50.0**, shared at the top by names such as Oregano and Crust Stone Oven Pizza. That ceiling is where tasting-menu and destination pizzerias sit.

Fast facts

The numbers that set the scale for this report:

10,000Records in the working dataset
0.00Median Price range min
50.0Highest observed Price range min
OreganoTop Name by Price range min

Data and method

The source is the TidyTuesday release from 2019-10-01 (pizza_datafiniti.csv). After cleaning, the analysis frame used here contains 10,000 records.

Price range minimum is the ranked metric across the chart stack. Medians are used because the distribution is heavily piled at low values. Charts are Plotly JSON with PNG fallbacks.

Price Floors Separate Everyday Shops From Destination Pies

Price range min by Name

At the top of the price-range-minimum ranking, Oregano and Crust Stone Oven Pizza both hit **50.0**. Just below, a cluster of shops — Papa’s Pizza, Apizza Scholls, Barbaro, La Montanara, Spin, Dough Pizzeria Napoletana — sits at **40.0**.

Those are not typical delivery-store floors. They mark the upper edge of how this dataset encodes premium positioning. Most of the 10,000-row catalog never approaches that band.

The Top Dozen Is a Premium Plateau

Crust Stone Oven Pizza leads at 50.0 — 40.0 marks the median among the top dozen

Crust Stone Oven Pizza appears among the leaders at **50.0**, and the median among the top dozen is **40.0**. Inside the expensive club, the story is less about a single outlier than about a shared premium floor.

That plateau is useful for citation: when someone asks what “expensive pizza” looks like in this file, the answer is a top-dozen median of 40 — not the global median of 0.

Almost Everything Sits in the Low Bins

Price range min Distribution

The distribution is extreme. Roughly **8,167** rows fall near the lowest price-range-minimum bin, with about **1,801** in the next populated band. Only a few dozen observations occupy the high bins near 40–50.

That is why the median of 0.00 is both technically correct and narratively incomplete. The mass of the market is cheap or zero-coded; the interesting cultural story lives in the thin right tail.

Premium Names Concentrate Quickly

Cumulative Price range min

The Pareto view of price-range minimum among leading names climbs steeply: the first five entries already hold roughly **43%** of the plotted aggregate, and the curve reaches the full total by the fifteenth name.

Concentration here is not about volume of pizzas sold — the file is not a sales ledger — but about how the high price-floor labels cluster among a short list of restaurant names.

Minimum and Maximum Ranges Move Together — Until They Don't

Price range min vs Price range max

Scatter points of price-range minimum against price-range maximum show the expected diagonal: shops with higher floors also tend to list higher ceilings, with some pairs reaching from 40 up toward 55.

The interesting cases are the mismatches — low floors with wide ceilings, or compressed ranges that signal a single-price menu. Those joints are where “pizza” stops being one product and becomes a price architecture.

What this file cannot tell you

Datafiniti restaurant scrapes are not audited menus. Zero price-range values may mean missing data rather than free food. Name collisions and chain-vs-independent ambiguity remain.

The extract cannot speak to ingredient costs, tip culture, or city-level rent. It measures listed price-range fields in a 2019 community snapshot.

What to take away

In 10,000 pizza rows, the market's mass sits at a median price-range minimum of 0.00, while a thin premium tier reaches 50.0.

Cite the top-dozen median of 40 when the question is about destination pizza; cite the global median when the question is about the typical listing. They describe different economies.

Sources

Data Science Learning Community. (2019). *TidyTuesday: All The Pizza*. https://raw.githubusercontent.com/rfordatascience/tidytuesday/main/data/2019/2019-10-01/pizza_datafiniti.csv

Which Ramen Brands Earn the Highest Ratings?

Instant ramen is a global industrial food that still gets reviewed like craft beer. Star ratings compress broth, noodle texture, spice, and packaging nostalgia into a single number that traveling eaters argue about for years.

A TidyTuesday working file of **3,180** records puts median stars at **3.75** and the highest observed stars at **5.00**. **Nongshim** leads the brand ranking by stars in the fact box, and **Japan** is the most common country label — a reminder that production geography and rating leadership are related but not identical.

A ceiling of 5.00 that many top brands share means the interesting story is often below the perfect scores: which countries sit above the median, and how review volume relates to praise.

Fast facts

The numbers that set the scale for this report:

3,180Records in the working dataset
3.75Median Stars

5.00Highest observed Stars
NongshimTop Brand by Stars
JapanMost common Country

Data and method

The source is the TidyTuesday release from 2019-06-04 (R for Data Science community). The working file contains 3,180 rows and 6 columns after cleaning — brand, variety, style, country, stars, and review number among the fields.

Medians handle a distribution piled against a hard 5.00 ceiling. Charts export as Plotly JSON with PNG fallbacks. Stars are community review scores from the source catalog, not lab nutrition assays.

Stars by brand

Stars by Brand

Brand rollups at the top of the chart crowd the **5.00** ceiling — names such as TTL, Tseng Noodles, A-Sha, Prima Taste, and related specialists appear in the perfect-score band. When the top of a rating system saturates, discrimination moves to consistency across varieties and to how many reviews support the mean.

Nongshim's fact-box lead by stars can coexist with a chart head full of smaller perfect-score brands. Aggregation choices — mean across varieties, review-weighted scores, or max — change who looks like the champion.

Who sits at the top

TTL leads at 5.00 — 5.00 marks the median among the top dozen

TTL leads at 5.00, and the median among the top dozen is also **5.00**. The elite band is a plateau, not a ladder. Perfect scores are common enough at the head that first place is a crowded room.

For readers, that plateau is a warning against over-reading tiny rank differences among five-star brands. The useful comparisons live in the middle of the distribution and across countries.

Stars by country

Stars by Country

Country box plots show whether star consensus is shared or contested across producing nations. Japan leads headcount; Malaysia, Thailand, South Korea, Singapore, and others compete on rating location rather than volume alone.

A country can dominate shelf presence without dominating the upper quartile of stars. The spread chart is where industrial scale and taste prestige separate.

Who beats the median — and who trails

Stars vs median by Country

Malaysia sits **0.50** above the median on the gap chart; **Thailand** trails by **0.25**. Those gaps look small until you remember the scale only runs to 5.00 and the median is already 3.75.

Half a star above the center is a meaningful national reputation effect in a compressed rating system. It is also small enough that a few legendary varieties can move a country's summary.

Stars and review number

Stars vs Review number

Plotting stars against review number shows clusters that averages erase. Some high-rated varieties have thin review support; some heavily reviewed products sit nearer the middle. Attention and admiration are related, not identical.

Review order fields in the source can also behave like identifiers as much as popularity counts depending on the release. The scatter's job is to keep readers from equating a perfect score with a broadly tested consensus.

What this file cannot tell you

Community-cleaned TidyTuesday snapshots are not live APIs. Missing values, spelling variants, and brand-name duplicates apply. Stars are not sales, and country labels are not always where a variety is eaten most.

Findings describe structural signals about ramen rating metadata — not a definitive world championship of instant noodles.

What to take away

Instant ramen ratings in this file center near 3.75 stars, with a crowded perfect-score ceiling at 5.00 and country gaps on the order of a quarter to half a star around the median.

The citable lesson is compression: when the top is a plateau of five-star brands, the meaningful geography is which countries sit above the middle and how much review support stands behind the praise.

Sources

Who Drinks the Most Alcohol Per Capita?

Global alcohol data is often collapsed into a single “drinks per person” headline. The TidyTuesday global alcohol extract keeps the beverage split intact — beer, spirits, and wine servings — across **193** country-level records. Beer is the metric that leads these charts, with a median of **76.0** servings and a recorded high of **376** in Namibia.

That spread is the story. A median of 76 and a maximum of 376 means the top of the beer distribution is not a gentle tail; it is a different consumption regime. Country rankings on beer alone will not match rankings on spirits, and the scatter at the end makes that explicit.

Research question

When alcohol data is separated by beverage, does a country's apparent drinking profile still look like one national habit, or does it split into distinct beer, spirits, and wine cultures? This report uses the 2018 TidyTuesday global alcohol extract to ask where beer volume concentrates and whether the beer leaderboard can stand in for total alcohol behavior.

The observational question is deliberately narrow: which countries sit in the high-beer tail, how far above the median they are, and whether beer intensity travels with spirits intensity. That focus keeps the charts from becoming a generic public-health ranking and instead treats beverage mix as a measurable cultural and market structure.

Fast facts

The numbers that set the scale for this report:

193Records in the working dataset

76.0Median Beer servings

376Highest observed Beer servings

NamibiaTop Country by Beer servings

Data and method

The source is the TidyTuesday release from 2018-06-26 (week13_alcohol_global.csv), maintained by the R for Data Science community. The working file used here contains 193 rows after cleaning.

Beer servings are the primary ranked metric in the chart stack. Medians are used where distributions skew. Charts are Plotly JSON with PNG fallbacks.

Beer Consumption Has a Steep Country Ladder

Beer servings by Country

Namibia leads the beer ranking at **376** servings, followed by the Czech Republic at **361**, Gabon at **347**, and Germany at **346**. Lithuania and Poland tie near **343**. Ireland sits lower in the top dozen at **313**.

These are not clustered scores around the global median of 76. The leaders occupy roughly four to five times the typical country's beer volume in this file — a reminder that beer culture is geographically concentrated, not evenly shared.

The country names matter because they point to different beverage systems. Namibia and Gabon place African markets in the same high-beer conversation as the Czech Republic, Germany, Lithuania, Poland, and Ireland — countries often associated with long-standing brewing industries, beer taxation regimes, and pub or lager cultures. A single “Europe drinks beer” shortcut would miss the African leaders; a single “income predicts beer” shortcut would miss the spread across very different economies.

The 2018 TidyTuesday file descends from World Health Organization alcohol indicators that summarize recorded servings rather than clinical outcomes. That distinction is important: a serving count can identify beverage concentration, but it does not measure binge patterns, liver disease, drunk-driving deaths, or unrecorded home production. The chart is strongest as a map of recorded beverage preference.

The Top Dozen Has Its Own Median

Namibia leads at 376 — 338 marks the median among the top dozen

Among the top twelve beer countries, Namibia's **376** still leads, while the median of that elite group is **338**. Even inside the high-beer club, the distribution has a ceiling that most of the world never approaches.

Reading only the global median would hide that club entirely. Reading only Namibia would hide how many Central European and other high-beer countries sit just behind it.

The top-dozen median of 338 is almost as revealing as the maximum. It says the leaders are not one outlier and eleven ordinary countries; they are a compact high-consumption band. Czech, German, Lithuanian, Polish, and Irish entries sit in a historically industrial beer belt, while Namibia and Gabon remind us that contemporary beer markets also reflect colonial trade routes, urban retail systems, and modern commercial breweries.

For public-health interpretation, this is where volume and harm begin to diverge. WHO and OECD alcohol reports usually emphasize liters of pure alcohol, age-standardized mortality, and sex-specific patterns because beverage-specific servings alone cannot compare the ethanol load of beer, wine, and spirits. This Artometrics cut keeps the serving metric because it asks a cultural-economy question first: where is beer, as a category, unusually dominant?

Most Countries Cluster Far Below the Peaks

Beer servings Distribution

The distribution chart piles most countries into lower beer-serving bins. Counts peak in the low double-digit and mid-range bins, while the extreme high bins that hold Namibia and the Czech Republic contain only a handful of countries.

That shape is why the median (76.0) is the right headline number: it describes the typical country, not the viral maximum.

The skew also changes how readers should interpret “per capita” comparisons. A country can sit below the median because beer is less available, because spirits or wine dominate the market, because alcohol consumption is lower overall, or because religious and legal rules suppress recorded sales. Those mechanisms are not interchangeable, even if they all land in the same low-beer bins.

The lower half of the distribution should therefore be read as a mixture of abstention, substitution, market structure, and reporting practice. In many Muslim-majority countries, for example, recorded beer

servings will reflect legal and religious constraints as much as consumer taste. In wine-producing countries, the same low beer value may simply mean the alcohol profile has moved to another column.

Concentration Rises Quickly Among the Leaders

Cumulative Beer servings

The Pareto curve shows how cumulative beer servings accumulate as countries are added from the top. The first five entries already account for about **37%** of the aggregate among the plotted leaders; by fifteen entries the curve reaches 100% of that restricted total.

Beer volume at the top is not democratic. A short list of countries carries a disproportionate share of the servings represented in this ranking view.

A Pareto view is useful because alcohol markets often look concentrated at multiple scales: a few beverage categories, a few dominant firms, and a few countries with unusually high recorded intake. The first five countries in this leader-only curve are not a global share of all alcohol; they are the head of the beer ranking. That narrower denominator is exactly why the curve is steep.

Concentration also affects policy comparison. If a small set of countries drives the visual range, then a dashboard built around maxima will exaggerate the distance between most countries. Median-centered views and beverage-specific scatterplots are better for separating ordinary national profiles from the high-consumption head.

Beer and Spirits Are Not the Same Map

Beer servings vs Spirit servings

Plotting beer servings against spirit servings across all 193 countries breaks the assumption that “heavy drinking” is one behavior. Some countries post high beer with moderate spirits; others invert the pattern.

The joint distribution is the corrective lens: beverage preference is cultural and economic, not a single alcohol intensity score. Namibia's beer lead does not automatically imply a spirits lead.

The scatter is the most important methodological guardrail in the report. Eastern European and post-Soviet alcohol research, for example, often emphasizes spirits because high-proof beverages can dominate harm profiles even when beer is not the leading column. Mediterranean comparisons often separate wine from both. A beer-only table cannot adjudicate those patterns.

That is why the article does not collapse the file into a total-servings index. Total alcohol may be the right measure for burden-of-disease work; beverage separation is the right measure for understanding culture, taxation, retail systems, and substitution. The same country can be a beer outlier, a spirits moderate, and a wine footnote.

What this file cannot tell you

Country-level serving estimates inherit survey, sales, and tourism distortions. Unrecorded alcohol, home production, and cross-border purchases can bias official figures. The 2018 TidyTuesday snapshot is not a live WHO dashboard.

Rankings on beer servings alone should not be quoted as total alcohol harm rankings. Spirits and wine columns exist precisely because the beverage mix differs.

What to take away

Beer consumption in this extract is a steep hierarchy: Namibia at 376 servings against a world median of 76. The top dozen forms its own high plateau around a median of 338.

The lasting insight is comparative: beer leaders and spirits leaders are not the same set of countries, and any serious alcohol map has to keep those axes separate.

Sources

Data Science Learning Community. (2018). *TidyTuesday: Alcohol Consumption*. https://raw.githubusercontent.com/rfordatascience/tidyuesday/main/data/2018/2018-06-26/week13_alcohol_global.csv

World Health Organization. (2018). *Global status report on alcohol and health 2018*. <https://www.who.int/publications/i/item/9789241565639>

World Health Organization. (2024). *Global status report on alcohol and health and treatment of substance use disorders*. <https://www.who.int/publications/i/item/9789240096745>

OECD. (2021). *Preventing Harmful Alcohol Use*. <https://doi.org/10.1787/6e4b4ffb-en>

Rehm, J., et al. (2009). Global burden of disease and injury and economic cost attributable to alcohol use and alcohol-use disorders. *The Lancet*, 373(9682), 2223–2233. [https://doi.org/10.1016/S0140-6736\(09\)60746-7](https://doi.org/10.1016/S0140-6736(09)60746-7)

Which Fast-Food Chains Pack the Most Calories Per Item?

Fast-food menus are marketing documents dressed as nutrition tables. The calorie counts attached to each item are among the few numbers that cut through the branding — and they vary wildly across chains that sell roughly the same categories of food.

A TidyTuesday snapshot of major U.S. chain menus puts **515** item records on the same footing. The median item sits at **490** calories. The ceiling in the file is **2,430** — a twenty-piece buttermilk crispy chicken tender tray that makes the median look almost restrained.

Taco Bell appears most often in the working table, but frequency of listing is not the same as caloric intensity. The story is which items and which chains sit above the center of the menu, and how fat tracks with calories when the branding falls away.

Fast facts

The numbers that set the scale for this report:

515Records in the working dataset

490Median Calories

2,430Highest observed Calories

20 piece Buttermilk Crispy CTop Item by Calories

Taco BellMost common Restaurant

Data and method

The source is the TidyTuesday release from 2018-09-04, maintained by the R for Data Science community. The working file contains 515 rows and 17 columns after merging the available CSV tables in the week folder — restaurant name, item, calories, total fat, and related nutrition fields.

Medians are preferred where distributions skew. Index-style fields and sequential IDs are excluded from metric selection. Charts ship as Plotly JSON with PNG fallbacks. The snapshot is a menu catalog for a fixed release window, not a live API of every current offering.

Calorie ceilings by item

Calories by Item

At the top of the item ranking, the **20 piece Buttermilk Crispy Chicken Tenders** lead at **2,430** calories. Related tender trays populate much of the upper tier: the 12-piece sits near **1,510**, while the 10-piece anchors a lower rung around **1,210**.

What looks like a single chicken category is really a portion machine. Doubling the piece count nearly doubles the calorie load. Burger King and Sonic doubles — Farmhouse King, Rodeo King, Super Sonic doubles — fill the same high band for a different reason: stacked patties, sauces, and bun mass rather than tray size.

Who sits at the top

20 piece Buttermilk Crispy Chicken Tenders leads at 2,430 — 1,315 marks the median among the top dozen

Among the top dozen calorie items, the median is **1,315**. That is more than two and a half times the dataset-wide median of 490. The head of the menu is not a gentle slope — it is a separate altitude.

Shareable trays and king-style sandwiches dominate the list. The implication for comparison is simple: chain averages without portion context will mis-rank brands that sell both modest sandwiches and extreme trays under the same logo.

How chains spread calories

Calories by Restaurant

Box plots by restaurant show whether a chain's calorie consensus is shared or contested. Some menus cluster tightly around the middle. Others stretch from light sides to extreme trays, producing wide interquartile ranges that averages erase.

Taco Bell's heavy representation in the file does not automatically place it at the caloric peak. Volume of SKUs and intensity of SKUs are different questions — and the distribution chart is where that split becomes visible.

Who beats the median — and who trails

Calories vs median by Restaurant

Relative to the dataset median, **Sonic** sits **80** calories above the center. **Chick-fil-A** trails by **100**. Burger King, Arby's, and McDonald's sit on the high side of the gap chart; Subway and Taco Bell sit below.

These gaps are menu-composition effects as much as health branding. A chain that lists more grilled items and fewer shareable trays will land below the median even if its signature sandwich is not light. The gap chart ranks catalogs, not single hero products.

Calories and fat move together

Calories vs Total fat

Plotting calories against total fat produces the expected positive slope — and the clusters that averages erase. Dense clumps of ordinary sandwiches sit in the middle; the extreme trays and stacked burgers occupy the upper-right corner where both axes spike together.

Fat is not a perfect proxy for calories, but in this menu file the two metrics rarely disagree about which items are extreme. When they do separate, it is usually because sugar-heavy drinks or desserts pull calories without matching fat.

What this file cannot tell you

Community-cleaned TidyTuesday snapshots are not live APIs. Missing values, spelling variants, and week-of-export coverage limits apply. Merged tables may fan out or duplicate rows when join keys are imperfect.

The file does not measure how often items are ordered, how portions are actually consumed, or how recipes have changed since the 2018 release. Treat the numbers as a structural map of listed menu calories — not a verdict on every current drive-through board.

What to take away

The median fast-food item in this catalog sits near 490 calories. The items that define the public argument about fast food often sit two to five times higher — trays and stacked sandwiches that are easy to order and hard to compare without portion context.

Chain gaps of eighty to one hundred calories around the median are real, but they are catalog effects. The clearest signal is the ceiling: once shareable chicken trays enter the frame, the conversation about a typical item and the conversation about the heaviest item are no longer the same conversation.

Sources

Which NYC Cuisines Score Highest on Health Inspections?

New York City restaurant inspections turn hygiene into a public grade. Behind the A/B/C stickers is a numeric score distribution that is far more skewed than the cheerful window cards suggest.

A TidyTuesday working extract of **100,000** records puts median score at **15.0** and the highest observed score at **156**. **NOUS LES AMIS RESTAURANT & B** leads the DBA ranking by score in the fact box, grade **A** is the most common label, and the year span in the file runs from a placeholder-early **1900** through **2018** — with the dense modern inspection era carrying most of the signal.

Higher scores here mean more violation points, not better food. That inversion is easy to forget once letter grades take over the conversation.

Fast facts

The numbers that set the scale for this report:

100,000Records in the working dataset

15.0Median Score

156Highest observed Score

NOUS LES AMIS RESTAURANT & BTop DbA by Score

1900–2018Year span covered in the file

AMost common Grade

Data and method

The source is the TidyTuesday release built from NYC restaurant inspection open data via the R for Data Science community. The working file contains 100,000 rows after sampling or assembly — DBA names, boroughs, scores, grades, and inspection dates among the fields.

Medians matter because a thin tail of high-violation restaurants pulls means upward. Charts export as Plotly JSON with PNG fallbacks. Grades are derived from scores under city rules; treating them as independent aesthetics misses the point system underneath.

How scores moved over time

Median Score Over Time

Median score over time tracks whether the typical inspected restaurant became cleaner, worse, or merely differently covered as the program matured. Flat stretches can mean stable compliance — or stable scoring practice.

Because early placeholder years can sit beside dense 2010s rows, the right-hand side of the timeline is where most interpretive weight belongs.

Who sits at the top of violation scores

The Slope Lounge and Restaurant leads at 152 — 122 marks the median among the top dozen

The Slope Lounge and Restaurant leads at **152** on the highlighted cut, with **122** as the median among the top dozen. Those are extreme violation totals relative to a file-wide median of 15.

The head of the score distribution is where enforcement drama lives. It is also where single bad inspections can define a DBA's appearance in a ranking even if later visits improve.

Scores by grade

Score by Grade

Box plots by grade show the numeric ranges that map onto A, B, C, and related labels. Grade A dominates headcount; the higher-letter grades occupy the upper score territory where violation points accumulate.

The chart makes the letter system's compression visible: many different numeric realities share an A, while the public argument concentrates on the rarer worse grades.

Who beats the median — and who trails

Score vs median by Grade

The gap chart ranks grades above or below the median score. Worse letter grades sit far above the median by construction; A sits at or below it. The geometry is the grading rule made visual.

What remains interesting is the within-grade spread on the previous chart: even A is a band, not a point, and policy fights often happen inside that band.

Scores and identifiers

Score vs Camis

Plotting score against CAMIS identifiers is mostly a diagnostic scatter: restaurant IDs are not a substantive X-axis, but the cloud still shows how scores disperse across the universe of establishments rather than clustering on a few famous names.

When substantive covariates are thin, an ID scatter is a honesty device — it refuses to invent a prettier relationship than the table supports.

What this file cannot tell you

Community-cleaned TidyTuesday snapshots are not live APIs. Missing values, spelling variants, and sampling to 100,000 rows apply. Scores are inspection outcomes, not Yelp taste ratings.

Findings describe structural signals about NYC inspection scoring in the file — not a complete health study, and not a recommendation about where to eat tonight without checking current grades.

What to take away

NYC restaurant inspections are a median-15 world with a long right tail past 150 violation points. Grade A is common; extreme scores are rare and loud.

The citable inversion remains: higher numbers are worse, letter grades compress a wide numeric band, and the public sticker is only the headline of a point system.

Sources