

# **The Identity Issue — Artometrics**

## **The Identity Issue**

Language, schools, regional identity — preference infrastructure made visible.

## **Where U.S. Student-Body Diversity Is Highest**

School-level composition data map where student-body diversity concentrates across U.S. local education agencies. The working file holds 27,944 records. The calibration metric is Multi, with a median of 2.38 — the center of the observed diversity scores in this extract.

The ceiling is stark: butteville union elementary leads at 85.3 Multi, while Texas (TX) is the most common state abbreviation in the file. High Multi at the top of the leaderboard and a modest median across the whole table are not contradictions. They are the shape of a heavily skewed education landscape.

## **Research question**

Where does student-body diversity concentrate when the unit of analysis is the local education agency rather than the nation as a whole? This report asks whether the Multi metric behaves like a broad national baseline or a long-tailed distribution in which a few LEAs and states sit far above the typical row.

The observational question is constrained to the source file: how stable is the median, which LEA names define the upper tail, how do state abbreviations structure the spread, and how does Multi relate to Variance? It does not treat one diversity metric as the full moral or demographic story of American education.

## **Fast facts**

The numbers that set the scale for this report:

27,944Records in the working dataset

2.38Median Multi

85.3Highest observed Multi

butteville union elementaryTop LEA NAME by Multi

TXMost common ST

## **Data and method**

The source is the TidyTuesday release from 2019-09-24 (R for Data Science community). The working file contains 27,944 rows and 15 columns after merging available tables in the week folder. Multi is the primary observed metric used for ranking and distribution charts; Variance appears alongside it in the scatter.

Medians are used because school-level diversity metrics skew. Index-style fields are excluded from metric selection. State codes (ST) and LEA names are categorical axes for comparing how Multi sits across geography and district identity.

## **A flat median hides how extreme the upper tail is**

### **Median Multi Over Time**

In the trend cut available in this file, median Multi sits at 2.38 in the opening period and 2.38 at the close — a flat line at the center of the distribution. Stability at the median does not mean stability everywhere; it means the typical school's Multi did not move in this snapshot window.

That flat center is the baseline against which the leaderboard should be read. The interesting action in school diversity here is not a national median march — it is how far individual LEAs sit above that center.

The flat median is especially important because U.S. school diversity is often narrated through landmark legal periods: *Brown v. Board of Education* in 1954, the busing and desegregation orders of the 1960s and 1970s, the release from court supervision in later decades, and contemporary residential segregation. Those histories are real, but this chart is not a long legal history; it is a snapshot-style metric view where the center does not move.

A stable center can coexist with sharp local change because public education is geographically partitioned. District boundaries, housing markets, charter growth, school assignment rules, and private-school exit can all alter individual LEAs without moving the national median much. The chart tells readers to look for distributional extremes rather than assume a single national trend.

## **Butteville Union Elementary leads a steep LEA ladder**

### **Top LEA NAME**

butteville union elementary leads at 85.3 Multi. Among the top dozen LEAs, the median is 52.0 — still many times the file-wide median of 2.38. The leaderboard is an upper-tail map of districts whose composition scores sit far from the typical school.

LEA names on this chart are the entities carrying the highest observed Multi values in the file. They are not a random sample of American schooling; they are the extreme cases the metric surfaces first.

Butteville Union Elementary is a named local education agency rather than a state system, which is exactly why the scale warning matters. A small district can produce an extreme composition metric if its enrollment mix is unusual, if denominators are small, or if categories split more evenly than in larger systems. The top of the ladder is a place to investigate, not a national template.

The top-dozen median of 52.0 confirms that the upper tail is not one data-entry spike. Several LEAs sit in a high-Multi band that is many times the file-wide center. In education data, that kind of gap often reflects the interaction of neighborhood demography, district boundaries, grade configuration, and reporting categories.

## **State boxes show how Multi spreads geographically**

### **Multi by ST**

Splitting Multi by state (ST) produces category boxes that reveal whether diversity scores are shared or contested across jurisdictions. Texas appears most often in the file, but frequency of rows is not the same as high Multi; volume and score must be read separately.

The state axis is useful because education policy and demography both vary geographically. The boxes show the within-state spread that a single national median erases.

Texas appearing most often is not surprising in a school file: it is a large state with many districts, fast-growing metropolitan regions, and substantial racial and ethnic diversity across Houston, Dallas-Fort Worth, San Antonio, Austin, the Rio Grande Valley, and rural counties. But row frequency mainly measures how much Texas appears in the extract, not whether every Texas LEA scores high on Multi.

State boxes are also institutional boxes. School finance rules, district consolidation, county governance, accountability systems, and enrollment reporting vary by state education agencies. A state-level spread therefore mixes demography with administrative design, which is why the box plot is more informative than a single state average.

## **Oklahoma clears the median; New Jersey trails**

### **Multi vs median by ST**

Relative to the median, OK sits 4.25 above; NJ trails by 0.81. That gap chart converts state boxes into signed distance from the file center — a quick way to see who systematically lands higher or lower on Multi.

These are structural offsets in the working extract, not causal claims about state policy. They locate where the distribution's weight sits once the national median is the reference line.

Oklahoma's positive gap and New Jersey's negative gap should be read against the metric definition and the units in the file. Oklahoma contains urban districts, rural districts, and tribal-nation contexts that can produce different student-body compositions across local agencies. New Jersey contains highly diverse metropolitan areas as well, but district boundaries and within-state sorting can produce lower values on this particular Multi summary.

The signed-distance format is useful because it resists a common error in education debate: assuming that states with diverse populations automatically have diverse schools. Residential segregation, municipal fragmentation, selective programs, and district borders can separate population diversity from school-level diversity.

## **Multi and Variance move together in clusters**

### **Multi vs Variance**

Plotting Multi against Variance shows joint structure that averages erase: schools and districts form clusters rather than a single smooth cloud. High Multi with different Variance profiles is a different composition story from moderate Multi with tight Variance.

The scatter's job is relational. It does not crown a winner; it shows which combinations of diversity level and dispersion appear together in the 27,944-row file.

Variance adds a second lens because two LEAs can have similar Multi values while distributing students across categories differently. One district may be broadly balanced across several groups; another may be split sharply between two groups; a third may have a small high-Multi subgroup inside a mostly homogeneous system. Multi alone cannot distinguish those patterns.

This is where school diversity connects to the measurement literature on segregation and exposure. Researchers often use dissimilarity, isolation, entropy, and exposure indices because no single measure captures both representation and distribution. The scatter does not replace those indices, but it shows why a one-column metric needs a dispersion companion.

## What this file cannot tell you

Community-cleaned TidyTuesday snapshots are not live federal education APIs. Missing values, spelling variants in LEA names, and week-of-export coverage limits apply. Merged tables may fan out or duplicate rows when join keys are imperfect.

Multi is the metric as defined in the source file. Findings describe this school-diversity extract — not a complete census of every U.S. school in every year.

## What to take away

The national median Multi of 2.38 is stable in the available trend window, while individual LEAs such as butteville union elementary reach 85.3. Diversity here is a long-tailed property: the center barely moves, and the extremes define the drama.

State gaps (Oklahoma above the median, New Jersey below) and the Multi–Variance scatter add geography and joint structure. Together they show where student-body diversity concentrates without pretending one number explains the whole map.

## Sources

Data Science Learning Community. (2019). *TidyTuesday: School Diversity*. [https://raw.githubusercontent.com/rfordatascience/tidyuesday/main/data/2019/2019-09-24/school\\_diversity.csv](https://raw.githubusercontent.com/rfordatascience/tidyuesday/main/data/2019/2019-09-24/school_diversity.csv)

Data Science Learning Community. (2019). *TidyTuesday School Diversity source folder*. <https://github.com/rfordatascience/tidyuesday/tree/main/data/2019/2019-09-24>

National Center for Education Statistics. *Common Core of Data*. <https://nces.ed.gov/ccd/>

Reardon, S. F., & Owens, A. (2014). 60 years after Brown: Trends and consequences of school segregation. *Annual Review of Sociology*, 40, 199–218. <https://doi.org/10.1146/annurev-soc-071913-043152>

## Editor's note

*Artometrics data report from the TidyTuesday research pipeline. Charts and aggregates are reproducible from the embedded exhibits and public source files.*

[View TidyTuesday source on GitHub](#)

---

# How Concentrated Is the World's Linguistic Diversity?

Linguistic diversity is often imagined as a smooth global gradient. Glottolog-style language inventories show something lumpier: documentation clusters by macroarea, and a long tail of languages appear once while a short head of names and identifiers repeats.

The TidyTuesday working file holds **8,612** records. **Africa** is the most common macroarea label, with **2,363** records in that bucket alone. The question is not only how many languages exist — it is where documentation effort, and therefore visibility, concentrates.

A catalog of languages is also a catalog of scholarly attention. Places with denser fieldwork traditions look richer in the table even when the underlying human diversity story is more even — or more endangered — than the row counts imply.

## Research question

What does a Glottolog-derived catalog reveal about the geography of linguistic documentation, not just the geography of language itself? This report asks where the 8,612-record extract concentrates by macroarea, where names and identifiers repeat, and how much of the visible language map is shaped by catalog infrastructure.

The question is observational because row counts are not speaker counts, vitality scores, or political claims. The aim is to distinguish language diversity as lived by communities from language visibility as preserved in databases, grammars, ISO-style identifiers, and scholarly citation systems.

## Fast facts

The numbers that set the scale for this report:

8,612Records in the working dataset

AfricaMost common Macroarea

## Data and method

The source is the TidyTuesday release from 2025-12-23 (R for Data Science community). The working file contains 8,612 rows and 9 columns after cleaning — language names, identifiers, macroarea labels, and related Glottolog metadata.

Because many fields are categorical rather than scored, the analysis leans on counts, concentration, and repetition rather than medians of a single numeric quality metric. Charts export as Plotly JSON with PNG fallbacks.

## A geographically uneven landscape

### Language documentation is geographically uneven

Africa dominates with **2,363** records in the primary macroarea cut. Other macroareas — Papunesia, Eurasia, the Americas, Australia — follow at lower counts. The main bucket carries much of the story; this field does not behave like a gentle long-tail split across dozens of equal regions.

Uneven documentation is not identical to uneven linguistic diversity, but it shapes what comparative science can see. Extinction risk, revitalization, and typology research all inherit the map of what has been recorded.

Africa's lead is plausible in linguistic terms: the continent contains large families such as Niger-Congo, Afro-Asiatic, Nilo-Saharan groupings as traditionally defined, and Khoisan-associated lineages, alongside intense contact zones and multilingual states. But the chart does not validate any particular family taxonomy; it records how the Glottolog-style file assigns macroarea metadata to rows.

Papunesia and Australia are especially important comparison cases because they contain many small language communities with high typological diversity and severe endangerment pressure. A lower row count than Africa does not mean lower scientific importance. It means the record is distributed unevenly across regions, fieldwork histories, and naming conventions.

## Who sits at the top of name repetition

### Repeated language names show documentation density

Repeated language names show documentation density inside the extract. **Fasu** appears among the recurring names in the file; the top dozen account for a visible share of all 8,612 rows even though most languages appear only once.

Name repetition can reflect dialect documentation, alternate glottonyms, or multiple database entries tied to related varieties. It is a signal to investigate, not an automatic claim that one language is more important.

Fasu is a Papuan language name associated with Papua New Guinea, and its recurrence is a reminder that catalog entries are not always one-to-one with what a lay reader calls a language. Varieties, dialect clusters, alternate names, bibliographic records, and genealogical nodes can all push the same surface name into repeated view.

This is where Glottolog's infrastructure differs from a census. A national census might ask who speaks a language at home; Glottolog asks how a language or lineage is classified and documented in the linguistic literature. The same community can be highly visible in one system and nearly invisible in another.

# Regional concentration

## Regional concentration shapes the language map

Africa is again the largest bucket with 2,363 records on the regional concentration chart. Category concentration shows where editorial and scientific attention should look first if the goal is to understand the file's center of gravity.

Mixed macroarea labels in the data — combinations such as Eurasia with Papunesia — remind readers that languages and contact zones do not always respect tidy continental borders.

The macroarea scheme is a practical database compromise. Eurasia contains Indo-European, Uralic, Turkic, Sino-Tibetan, Japonic, Koreanic, Dravidian, and many other lineages; the Americas contain hundreds of Indigenous languages with colonial-era documentation gaps; Papunesia compresses a spectacularly dense set of Papuan and Austronesian histories into one label. The labels help chart structure, but they are not theory in themselves.

Regional concentration also encodes academic history. Missionary grammars, colonial surveys, national language academies, SIL International work, university field programs, and contemporary community-led documentation all leave traces in catalogs. The chart is therefore a map of evidence production as much as a map of human speech.

## Most languages appear once

### Most languages appear once while a small head repeats

Most name entities appear only once; a small head revisits repeatedly. That power-law shape is typical of catalog-style tables: a few repeatedly documented varieties, and a long inventory of singleton entries.

For endangered-language work, the singleton majority is the point. Visibility in a database is already a scarce resource; languages that appear once may still be under-described relative to their social importance.

A singleton entry can represent a language with thousands of speakers, a severely endangered variety with only elders remaining, or a poorly documented name whose classification remains uncertain. Without speaker counts, vitality levels, and bibliographic depth, the frequency chart should be read as documentation structure, not community scale.

UNESCO's language vitality framework and Ethnologue-style speaker estimates answer different questions than Glottolog identifiers. The ArtoMetrics chart is valuable precisely because it shows the catalog layer that those other systems need to join against before scholars can ask about endangerment, education policy, or intergenerational transmission.

## Identifiers as metadata

### Identifier fields are metadata rather than a reader-facing thesis

Glottolog-style IDs such as **fasu1242** are the most repeated identifiers in the extract. Secondary dimensions add context when the primary table has no numeric score column — they are join keys for typology databases, not a popularity contest.

Reading IDs as if they were rankings misunderstands the infrastructure. The useful claim is that documentation systems create machine-readable handles that make some languages easier to cite, map, and study than others.

Identifier fields are the quiet machinery of comparative linguistics. Glottocodes such as **fasu1242**, ISO 639-3 codes, and family IDs let researchers connect grammars, lexicons, sound inventories, phylogenies, and geographic coordinates without relying on unstable spelling variants. The identifier is not the language; it is the database handle that lets the language remain findable.

That infrastructure has real consequences. A language with a stable identifier can be linked into typological databases such as WALS, Glottolog, PHOIBLE, or Grambank more easily than a language whose names are fragmented across archives. Machine readability becomes a form of scholarly visibility.

# What this file cannot tell you

Community-cleaned TidyTuesday snapshots are not live APIs. Missing values, spelling variants, and release coverage limits apply. Row counts are not speaker counts, and macroarea labels are not political maps.

The file cannot by itself measure vitality, endangerment, or intergenerational transmission. It measures documentation structure — a necessary but incomplete layer of the global language story.

## What to take away

World linguistic diversity, as seen through this Glottolog-derived extract, is geographically lumpy: Africa leads the macroarea counts at 2,363 of 8,612 records, while most language names appear once.

The citable lesson is about visibility. What scholars can count depends on where fieldwork and cataloging have been dense — and that density is itself a cultural and scientific fact.

## Sources

Data Science Learning Community. (2025). *TidyTuesday: Glottolog / language data*. <https://github.com/rfordatascience/tidyuesday/tree/main/data/2025/2025-12-23>

Hammarström, H., Forkel, R., Haspelmath, M., & Bank, S. *Glottolog*. Max Planck Institute for Evolutionary Anthropology. <https://glottolog.org/>

UNESCO. *Atlas of the World's Languages in Danger*. <https://www.unesco.org/languages-atlas/>

Dryer, M. S., & Haspelmath, M. (Eds.). *The World Atlas of Language Structures Online*. <https://wals.info/>

Skirgård, H., et al. (2023). Grambank reveals the importance of genealogical constraints on linguistic diversity and highlights the impact of language loss. *Science Advances*, 9(16). <https://doi.org/10.1126/sciadv.adg6175>

---

# How American Cities Build Distinct Sports Identities

A team is never just a team. It is attached to a region, a media market, a league structure, and a local memory system. That context changes what the same number means.

this piece turns the sports canon toward geography. The goal is to show how regions become recognizable under the data microscope.

Start with the scale: **10** — Large markets/cities compared; and **7** — Regions grouped for title-density analysis.

## Fast facts

The numbers that set the scale for this report:

10 Large markets/cities compared

7 Regions grouped for title-density analysis

6 Identity archetypes used in the map

94 Toronto drought/pain index in this editorial model

43B Approximate combined New York major-team value

## Data and method

The charts use public franchise histories, rounded valuation summaries, and editorial identity indices. The indices are not official statistics; they are structured prompts for comparing markets.

A cultural economist would ask how facilities, media density, migration, ownership, and league rules shape local identity. this piece is the first map layer for that conversation.

## Regional Title Density

### Championship memory clusters by region

A region's sports identity is not only current fandom. It is the sediment of old championships, old rivalries, and inherited expectations.

The Northeast and Great Lakes have unusually deep archives because their leagues, cities, and media systems matured early.

## **Market Stack**

### **Deep markets create sports media ecosystems rather than single-team identities**

New York and Los Angeles are not simply large markets. They are sports ecosystems where multiple teams compete for attention, mythology, and local legitimacy.

That makes cross-report comparison useful: a Knicks drought and a Lakers title live in the same media bloodstream.

## **Sport Dna**

### **Regions specialize in different sports identities**

Texas reads as football ritual. California leans toward basketball and glamour. The Northeast keeps baseball memory unusually alive.

The point is not stereotype; it is signal. Regional sports culture gives the same win-loss record a different accent.

## **City Pain versus Success**

### **A city can be successful overall and still carry a famous wound**

Multi-team cities do not experience success cleanly. A Boston fan can live near constant banners while a specific rival city experiences one drought as civic identity.

Pain is often team-specific, but markets distribute and amplify it.

# Identity Types

## Sports cities are portfolios of myths

Some places sell archive. Some sell stars. Some sell football Sundays. Some sell hunger because they have not yet received the validating title.

This is the bridge to future city bioeconomics reports: sports data is one part of a place's cultural fingerprint.

## What to take away

The main finding is that sports identity is regional before it is rational. Market size, title archives, sport specialization, and drought all shape the meaning of a fan base.

This gives the next batch a new tool: every team profile can now ask not only how good the team is, but what kind of place produced that version of good, bad, rich, cursed, or glamorous.

## Sources

Reference franchise histories from Baseball Reference, Basketball Reference, Pro Football Reference, and Hockey Reference.

Forbes. Professional sports franchise valuation lists.

U.S. Census regional framing, used only as a broad geography reference.

Public attendance, championship, and market summaries from league history pages.

## Editor's note

Regional and identity indices are editorial approximations designed to organize cross-report thinking. They should be refined as more city/team reports are added.

---

# How Population Scale Relates to Prosperity Across Countries

Population is the blunt instrument of international comparison: China and India alone redraw every map. The TidyTuesday CIA World Factbook extract used here holds **259** records with a median population of **5,220,371** and a maximum of **1,355,692,576** for China.

Mean population in the distribution annotation (~32 million) sits far above the median — the classic right skew of a world with a few demographic giants and a long list of small states.

## Fast facts

The numbers that set the scale for this report:

259Records in the working dataset  
5,220,371Median Population  
1,355,692,576Highest observed Population  
ChinaTop Country by Population

## Data and method

The source is the TidyTuesday release from 2024-10-22 (cia\_factbook.csv). After cleaning, 259 rows remain.

Population is the primary ranked metric; birth rate appears in the joint scatter. Charts are Plotly JSON with PNG fallbacks.

## A Handful of Giants Dominate Headcount

### Population by Country

China leads at about **1.36 billion**, India at about **1.24 billion**, then the European Union aggregate near **511 million**, the United States near **319 million**, Indonesia near **254 million**, and Brazil near **203 million**.

These are not peers of the median country at 5.2 million. They are a different scale class — and they drive every concentration statistic that follows.

## **Even the Top Dozen Has a Steep Internal Drop**

**China leads at 1,355,692,576 — 199,415,584 marks the median among the top dozen**

China still leads at 1,355,692,576, while the median among the top dozen is about 199,415,584 — less than one-sixth of China's count.

“Large country” is a pyramid, not a plateau. Second-tier giants are already far smaller than the demographic summit.

## **Most Countries Sit in the Small-Population Mass**

**Median 5,220,371 vs mean 32,294,361 — the shape is right-skewed**

The histogram's first bin holds the overwhelming majority of entries (about **218** of them), with only a handful of observations in the billionaire and near-billionaire bins.

That shape is why the median (5.2 million) is the honest typical-country number, and why the mean (~32 million) is pulled relentlessly upward.

## **The Top Five Already Hold Most of the People**

**The top 5 country entries account for 72% of the aggregate population**

The Pareto curve shows the top five country entries accounting for about **72%** of aggregate population in this ranking view. By fifteen entries the curve approaches the full plotted total.

Global population is not evenly distributed across flags. It is concentrated in a short list of demographic powers.

## **High Birth Rates Cluster Among Smaller Populations**

### **Population vs Birth rate**

Population versus birth rate places some of the highest birth rates (above ~40 per 1,000) among countries well below the mega-population tier, while several large countries sit at lower birth rates.

The joint distribution warns against assuming that demographic giants are also fertility leaders. Scale and birth rate are different stories in the Factbook.

## **What this file cannot tell you**

Factbook figures are estimates with uneven vintage and methodology across entities. Including the European Union alongside countries double-counts if users sum naively. Small territories and disputed entities add definitional noise.

Population totals are not prosperity scores. Pairing with birth rate is descriptive, not a development ranking.

## **What to take away**

In 259 Factbook rows, the median country has about 5.2 million people while China exceeds 1.35 billion. The top five entries hold roughly 72% of the aggregate in the concentration view.

Cite the median for a typical state; cite China/India and the Pareto share when the question is planetary concentration.

## Sources

Data Science Learning Community. (2024). *TidyTuesday: CIA World Factbook*. [https://raw.githubusercontent.com/rfordatascience/tidytuesday/main/data/2024/2024-10-22/cia\\_factbook.csv](https://raw.githubusercontent.com/rfordatascience/tidytuesday/main/data/2024/2024-10-22/cia_factbook.csv)

---

# Which UN Votes Split the Chamber Most Sharply?

United Nations roll-call votes are a diplomacy archive written as a contingency table. This working extract holds 100,000 records. The most common vote value is yes — not a moral claim, a frequency claim. The charts ask where that yes-mass concentrates, which countries appear most often, and how long-tailed country participation becomes.

Voting records identify resolutions that divided member states, but the first job of this file is simpler: map the landscape of yes, no, and abstain, then see which member names and country codes recur in a 100,000-row slice of UN voting history.

## Research question

What does a 100,000-row extract of United Nations roll-call votes reveal before any ideological coalition model is fitted? This report asks how vote categories are distributed, which member-state names recur most often, and whether country-code metadata confirms or complicates the entity picture.

The question is intentionally descriptive. A yes-heavy table does not mean the General Assembly lacks conflict; it means this sample's recorded positions are concentrated in one category. The analysis therefore starts with vote mass and country recurrence before moving toward more interpretive claims about alignment.

## Fast facts

The numbers that set the scale for this report:

100,000Records in the working dataset

yesMost common Vote

## **Data and method**

The source is the TidyTuesday release from 2021-03-23 (R for Data Science community). The working file contains 100,000 rows and 4 columns after merging available tables in the week folder. Vote is the primary categorical field; country and country code provide entity axes; frequency charts summarize repetition.

There is no single numeric “score” column driving a median leaderboard here. The analysis is categorical: counts, shares, and recurrence. Index-style fields stay out of metric selection.

## **Yes votes dominate the sampled UN record**

### **Yes votes dominate the sampled UN record**

Yes dominates with 79,663 records. In a 100,000-row extract, that is the center of gravity: most recorded positions in this slice are affirmative. Abstain and no exist, but they do not match yes for volume.

Dominance here is a frequency fact about the file, not proof that the General Assembly always agrees. It tells you where the mass of the table sits before any conflict analysis begins.

The United Nations General Assembly often votes on resolutions that have already passed through committees, negotiations, amendments, and diplomatic bargaining. A large yes category can therefore reflect agenda-setting and consensus-building before the roll call, not an absence of disagreement in world politics. The public vote is the final visible record of a process that began earlier.

UN voting scholars usually move beyond raw categories into ideal-point models, affinity scores, or issue-area subsets because the meaning of a yes vote depends on the resolution. A yes on decolonization, a yes on nuclear disarmament, and a yes on a budgetary item are all affirmative votes, but they do not carry the same geopolitical signal.

## **Brazil appears most often among country names**

### **Brazil appears most often in this extract**

Brazil appears 747 times — the most recurring country name in the file. The top dozen countries account for a visible share of all 100,000 rows, which is what a recurring-member chart is for: showing who is densely present in the extract.

High appearance counts can reflect coverage, membership longevity, or how the sample was built. They are presence metrics, not win rates.

Brazil is a useful leader because it has been a United Nations member since 1945 and is a frequent actor in General Assembly voting records. The country's foreign-policy tradition includes visible positions on development, non-intervention, South-South cooperation, and Security Council reform. But the 747 appearances shown here are still a row-count fact, not a claim that Brazil is uniquely influential in every issue area.

Country recurrence is also shaped by institutional continuity. Newer UN members, dissolved states, renamed states, and countries with interrupted or transformed membership histories will appear differently from long-standing members. The Soviet Union/Russia, Yugoslavia/successor states, Czechoslovakia/Czechia and Slovakia, and country-name standardization issues are exactly why entity handling matters in UN vote data.

## **Vote categories restate the file's center of gravity**

### **Vote categories reveal the file's center of gravity**

Yes is again the largest bucket with 79,663 records when vote categories are plotted as a bar landscape. Repeating the category cut is useful: it confirms that the headline frequency is not an artifact of a single chart type.

Category concentration shows where editorial attention should start — with the yes majority — while still leaving room to inspect the smaller no and abstain bands.

Abstentions deserve special caution because they can mean several things: neutrality, strategic ambiguity, alliance pressure, domestic constraint, protest against wording, or unwillingness to oppose a popular resolution directly. In Cold War voting studies, abstention is often treated as analytically distinct from both yes and no because it can carry diplomatic information without joining either pole.

No votes are smaller in this extract, but they are often more diagnostic. A rare no can identify a strongly contested resolution, a bloc split, or a state protecting a core interest. The category chart therefore gives the denominator; substantive research still needs resolution identifiers and issue tags to interpret the minority categories.

## **Country appearances follow a long-tail pattern**

### **Country appearances follow a long-tail pattern**

Most country entities appear only once in the frequency view; a small head revisits repeatedly. That power-law shape is typical of catalog-style international tables: a few members are densely observed, many are sparse.

Long tails warn against over-reading rare countries as if they had the same sample depth as Brazil-scale recurrence.

The long tail can reflect countries that entered the United Nations later, states that changed names, country labels split across historical entities, or sample construction that does not preserve every roll call for every member equally. Tuvalu, South Sudan, Czechia, and North Macedonia do not have the same historical runway as Brazil, India, France, or the United States.

For comparative work, this means country-level rates need exposure denominators. A country observed 747 times can sustain a more stable voting-profile estimate than a country observed only a handful of times. Presence depth is a statistical precondition for diplomatic interpretation.

# Country codes add a second identity layer

## Country codes add metadata rather than a new thesis

BR is the most repeated country code in the extract — the code-level echo of Brazil's name-level lead. Secondary dimensions like codes add join keys and metadata when the primary table has no numeric score column.

Read codes as infrastructure for linking, not as a rival thesis to the vote-share story. They confirm entity identity across the 100,000-row sample.

Country codes matter because names are historically unstable. “Congo,” “Democratic Republic of the Congo,” “Congo, Democratic Republic,” “Russia,” “Russian Federation,” and former state names can fracture a table if the join key is text alone. Codes make it possible to connect UN votes to region, income group, alliance, population, or trade data without hand-cleaning every label.

The BR code leading alongside Brazil's name indicates that the entity layer is internally coherent for the top case. That does not guarantee the whole file is perfectly standardized, but it shows why the extra chart exists: metadata fields are the connective tissue that allow UN voting records to become analyzable political data.

## What this file cannot tell you

Community-cleaned TidyTuesday snapshots are not the full UN voting API. Missing values, country-name variants, and week-of-export coverage limits apply. A 100,000-row extract may be a sample or a capped export rather than every roll call ever taken.

Findings describe this UN votes file — structural signals about vote categories and country recurrence — not a complete theory of international coalitions.

## What to take away

In this extract, yes is the majority language of recorded votes (79,663 of 100,000), Brazil is the most recurring country name (747 appearances), and BR leads among country codes.

The long tail of country appearances means presence is uneven. Start with the yes majority and the dense head of member recurrence; treat sparse countries as thin evidence until coverage is checked.

## Sources

Data Science Learning Community. (2021). *TidyTuesday: UN Votes*. <https://raw.githubusercontent.com/rfordatascience/tidytuesday/main/data/2021/2021-03-23/unvotes.csv>

Data Science Learning Community. (2021). *TidyTuesday UN Votes source folder*. <https://github.com/rfordatascience/tidytuesday/tree/main/data/2021/2021-03-23>

Erik Voeten. *United Nations General Assembly Voting Data*. Harvard Dataverse. <https://doi.org/10.7910/DVN/LEJUQZ>

Bailey, M. A., Strezhnev, A., & Voeten, E. (2017). Estimating dynamic state preferences from United Nations voting data. *Journal of Conflict Resolution*, 61(2), 430–456. <https://doi.org/10.1177/0022002715595700>

United Nations Digital Library. *Voting data and General Assembly records*. <https://digitallibrary.un.org/>

United Nations. *Member States*. <https://www.un.org/en/about-us/member-states>

## Editor's note

*Artometrics data report from the TidyTuesday research pipeline. Charts and aggregates are reproducible from the embedded exhibits and public source files.*

[View TidyTuesday source on GitHub](#)