

The Power Issue — Artometrics

The Power Issue

Leaders, silence, access brands, and the economics of when institutions stop protecting.

Julius Caesar and the Cost of Managing Power

Economics comes from the Greek *oikonomia* — the management of the house. Psychonomics asks the prior question: who is managing the soul of the house, and what does their inner architecture cost everyone living inside it? Julius Caesar is the founding case study. His policies were not abstract theories. They were his biography, scaled.

Caesar was born in 100 BCE into the Subura — Rome's roughest district — to a patrician family with a noble name and thin wallets. His father died when Caesar was fifteen, leaving him head of household with no resources and no patron. By the time he was twenty-five, Cilician pirates had captured his ship, set a ransom of twenty talents, and watched Caesar laugh, volunteer to pay fifty, spend thirty-eight days as their captive writing poetry and calling them illiterates, then return after his release with a fleet he raised himself and crucify every one of them — exactly as promised. This is not a fun anecdote. It is the earliest behavioral evidence of his core operating theory: the self is the institution. The pattern never changed. The scale did.

By 44 BCE, Caesar had crossed the Rubicon, won a civil war, reformed the debt system, redistributed land to 80,000 veterans and urban poor, rebuilt the calendar, and declared himself dictator for life. He was then stabbed twenty-three times by sixty senators — led by Brutus and Cassius, two men he had personally pardoned, appointed to the praetor-

ship, and promised future consulships. The house he built outlasted him by five centuries. The soul that built it was load-bearing in a way no institution can survive.

Fast facts

The numbers that set the scale for this report:

100 BCE Caesar's birth year — born into Rome's most turbulent political century, to a patrician family living in the Subura tenements

38 days Held captive by Cilician pirates at age 25 — he laughed at their ransom demand, volunteered to pay more, then crucified every one of them after his release

0.95 Caesar's Need for Power score under the Hermann (1999) Leadership Trait Analysis framework — near the ceiling of any world leader ever profiled

80,000 Veterans and urban poor resettled in land colonies during Caesar's dictatorship — his economic reforms addressed real structural problems and built his client base simultaneously

194 yrs How long after Caesar's assassination Rome reached its peak GDP per capita — he authored the Pax Romana's prosperity without living a day of it

23 Stab wounds on the Ides of March, 44 BCE — inflicted by 60 senators, led by two men Caesar had personally pardoned, appointed, and promised future consulships

Data and method

this piece applies the Leadership Trait Analysis (LTA) framework, developed by political psychologist Margaret Hermann (1999) and used operationally by the CIA's Center for the Analysis of Political Behavior. LTA scores leaders on seven psychological variables by analyzing the content of their speeches and writings: Belief in ability to control events (B), Conceptual complexity (CC), Distrust of others (D), In-group bias (IB), Need for power (P), Need for affiliation (A), and Task orientation (T). Scores are normalized against a reference group of 100+ world leaders profiled between 1945–2010, on a 0–1.0 scale.

Applying LTA to an ancient figure is a legitimate extension of the method — published scholarship has done so for historical leaders using the same coding rubric (see: Dyson, 2006, *Political Psychology* ; Post, 2004, *Leaders and Their Followers in a Dangerous World*). For Caesar, scores are coded from three primary sources: the *Commentarii de Bello Galli-*

co (58–50 BCE, self-authored), Plutarch's Parallel Lives: Caesar (c. AD 100), and Suetonius's Divus Julius (c. AD 121). Comparative scores for Mao, Xi, and Thatcher are drawn from Wang & Guo (2019) and the Hermann (1999) reference baseline.

The Roman GDP per capita series is drawn from Maddison (2007), with revisions by Scheidel & Friesen (2009) and Lo Cascio & Malanima (2011). All values are in 1990 Geary-Khamis (international) dollars — the standard comparative unit for pre-industrial economies. The Maddison Project's earliest reliable Roman anchor is 14 AD; estimates before that date carry wider uncertainty intervals and should be read as directional rather than precise. The uncertainty band on Chart 2 reflects the low-to-high range across published scholarly estimates.

Leadership Trait Analysis

Caesar's LTA profile has a distinctive asymmetric shape that makes the psychonomic argument visible before you read a word

quartz_off_screen 2 Caesar's LTA profile has a distinctive asymmetric shape that makes the psychonomic argument visible before you read a word. Two variables spike toward the outer ring — Need for Power ($P = 0.95$) and Belief in Ability to Control Events ($B = 0.92$). One variable collapses toward the center — Distrust of Others ($D = 0.31$). That triangle is the architecture of his career. He believed, with near-total conviction, that he could personally determine outcomes and that he should be the one doing so. He also extended trust readily and withdrew it slowly. The combination is extraordinarily effective — until the people you trust decide the same values you embody (the right of the strongest to rule) apply to removing you.

Caesar exceeds the average world leader baseline on every variable except Distrust, where he falls dramatically below average (0.31 vs. 0.48). He was not a more extreme version of a normal leader. He was a normal leader in every domain, but with pathologically low distrust and pathologically high power-need running simultaneously. Mao's comparison is instructive: Mao scores 0.74 on Distrust and 0.81 on In-Group Bias — the two variables that explain the Cultural Revolution. Both men scored near-ceiling on Need for Power. Mao purged. Caesar pardoned. Mao's successors had to dismantle his cult brick by brick. Caesar's successors built an empire on his name.

The single most important number on this chart is $D = 0.31$. It is not a character flaw and it is not a strength. It is a network architecture decision. A leader who extends trust this freely, at this level of power concentration, is building a system with no redundancy. Every node that fails hits the center directly. Chart 3 shows what that looks like when three nodes fail at once.

The Economic Fingerprint

The most important feature of this chart is not Caesar's era — it is what comes after it

The most important feature of this chart is not Caesar's era — it is what comes after it. Roman GDP per capita during the late Republic sat at approximately \$440–\$490 (1990 international dollars), barely clearing the subsistence floor of \$400 (Maddison, 2007). This was not a prosperous society on a per-person basis. The Scheidel-Friesen (2009) income distribution model places the top 1.5% of Romans capturing roughly 20% of all income, with the vast majority living near subsistence. Grain supply chains failed regularly. Debt bondage was common. The fiscal machinery of the Republic was increasingly captured by the same factions fighting civil wars for control of it.

Caesar's dictatorship lasted less than five years — 49 to 44 BCE. His reforms held after his death, but the GDP data shows no immediate uplift. The economy barely moves during Caesar's lifetime. It moves under Augustus, who took Caesar's policies and embedded them in institutions designed to outlive any individual: the principate, the imperial cult, the professional army with fixed terms and state pensions. From Augustus's consolidation (27 BCE) to Trajan (100 CE), Roman GDP per capita rose roughly 30% — from ~\$480 to ~\$630. The Scheidel-Friesen peak of ~\$680 at 150 CE, Rome at its maximum economic output, came 194 years after Caesar's assassination.

This is the psychonomic paradox the chart makes numerical. Caesar's theory of power was that the self could substitute for institutional architecture — that personal authority, sufficiently magnanimous and forceful, was durable enough to hold the gains he made. Augustus understood what Caesar did not. He built the house so that no single soul was load-bearing. Caesar authored the prosperity line. Augustus inhab-

ited it. The distance between the assassination marker and the GDP peak is not a tragedy of timing. It is the measurement of the difference between the two men's theories of power.

The Patronage Network

The network chart makes the psychonomic argument structural

The network chart makes the psychonomic argument structural. Every edge connects to Caesar at the center. There is no distributed authority, no secondary loyalty cluster, no institution that holds if the center fails. This is not how republics are supposed to work — it is how patronage networks work when one person has absorbed enough of the Republic into themselves that the two become indistinguishable. Caesar did not intend to build a fragile system. He built the only system his operating theory allowed: one where the self is the institution, and the institution is therefore exactly as durable as the self.

The three dashed red edges are the chart's payload. Decimus Brutus was in Caesar's inner military circle — the general he trusted as much as Antony. Marcus Brutus had fought against Caesar in the civil war, been pardoned, been given the praetorship, and been promised a future consulship. Cassius had done the same. All three were among the highest-trust, highest-centrality nodes outside Caesar's immediate family. Decimus Brutus was the man who physically escorted Caesar to the Senate on the morning of March 15, 44 BCE, after Caesar had considered staying home. The network assigned maximum proximity to the men who carried the knives.

Caesar's D score of 0.31 on Chart 1 was not naivety. It was the logical extension of his core belief: that he could convert anyone through the experience of his own magnanimity. He was, in the vast majority of cases, correct. He converted pirates into cautionary tales, Gallic chiefs into clients, and grain riots into political capital. The failure mode is right there in the chart. Trust everyone, control everything works at every scale — until three nodes decide simultaneously that the values Caesar embodied apply to removing him. There was no redundancy to absorb it.

What this file cannot tell you

This report applies the Leadership Trait Analysis framework — developed for living leaders with large, verifiable speech corpora — to an ancient figure whose primary text is self-authored political propaganda. The *Commentarii de Bello Gallico* was written by Caesar during the campaigns it describes, distributed to a Roman political audience, and never intended as an impartial record. LTA scores derived from it likely underestimate Distrust (since trust is the policy Caesar publicly advertised) and may overestimate Conceptual Complexity (since nuance served his rhetorical interests). These biases should be read as a floor and ceiling, not as distortions that invalidate the analysis. The LTA scores in this report reflect the psychological signature Caesar projected and acted upon — which is precisely what Psychonomics requires.

The Roman GDP series carries wide uncertainty intervals before 14 AD, the earliest reliable Maddison Project anchor. All pre-Augustus estimates are reconstructed from indirect evidence — grain prices, wage records from Egyptian papyri, archaeological trade volume proxies, and demographic models. Scheidel himself cautions that ancient GDP estimates “can only be regarded as rough approximations” given the paucity of surviving data. The GDP chart in this report should be read as a directional argument — the shape of the trajectory and its inflection points relative to Caesar’s era — not as a claim to precision at the level of individual data points.

The patronage network is qualitatively constructed from primary sources, not computationally derived from a social network dataset. Edge types and trust weights reflect documented historical events (appointments, pardons, correspondence) and the scholarly consensus drawn from Plutarch, Suetonius, and Strauss (2015). The network cannot capture relationships we don’t have records of, and it represents a snapshot of Caesar’s final years rather than the evolution of his alliances over time. A complete Psychonomics database — with multiple figures and longitudinal network data — would allow more rigorous comparative claims. That database is the long-run goal of this series.

What to take away

Julius Caesar is the founding case study for Psychonomics because his life is a closed experiment with a clean result. We have his biography in granular detail, his own writings, three independent ancient ac-

counts, and 2,000 years of institutional outcomes to measure against his decisions. The experiment result is unambiguous: the man who genuinely reformed Roman debt law, redistributed land to 80,000 veterans and poor, stabilized grain supply, rebuilt the calendar, and laid the administrative foundation of a 500-year empire never lived in the house he built. He was killed by the men he trusted most, using the values he had taught them, nineteen months after declaring himself dictator for life.

The three charts tell the same story from three angles. The LTA radar shows the psychological architecture — the specific combination of near-ceiling power-need and near-floor distrust that made Caesar simultaneously extraordinary and structurally exposed. The GDP series shows the economic gap — Caesar in the flat zone of the late Republic, Augustus and his successors riding the 30% growth curve built on Caesar's institutional reforms. The patronage network shows the mechanism — the dashed red edges pointing directly at the center, maximum trust assigned to the maximum-risk nodes, no redundancy, no institutional buffer between Caesar's confidence and its consequences.

Economics tracks how houses are managed. Psychonomics asks who is managing the soul of the house — and what it costs. In Caesar's case, the cost was paid by everyone who came after him, and they built something extraordinary with it. The soul that managed the house collapsed it in the process of perfecting it. The house that replaced it was built to survive without him, and it did — for fifteen centuries in the East, five in the West, and indefinitely in the calendar we still use today.

Sources

Bolt, J., & van Zanden, J.L. (2020). *Maddison Project Database, version 2020*. University of Groningen.

Caesar, J. (58–50 BCE). *Commentarii de Bello Gallico*. Trans. W.A. McDevitte & W.S. Bohn. Project Gutenberg.

Cicero, M.T. (46 BCE). *Brutus*.

Dyson, S.B. (2006). Personality and foreign policy: Tony Blair's Iraq decisions. *Foreign Policy Analysis*, 2(3), 289–306.

Goldsworthy, A. (2006). *Caesar: Life of a Colossus*. Yale University Press.

- Hermann, M.G. (1999). *Assessing Leadership Style: A Trait Analysis*. Social Science Automation.
- Lo Cascio, E., & Malanima, P. (2011). Ancient and pre-modern economies: GDP in the Roman Empire and early modern Europe. Paper presented at the Conference on Quantifying Long Run Economic Development, Venice, March 22–24.
- Maddison, A. (2007). *Contours of the World Economy 1–2030 AD*. Oxford University Press.
- Milanovic, B., Lindert, P.H., & Williamson, J.G. (2011). Pre-industrial inequality. *The Economic Journal*, 121(551), 255–272.
- Pliny the Elder. *Naturalis Historia*, 7.54.
- Plutarch. (c. AD 100). *Parallel Lives: Caesar*. Trans. Robin Seager. Penelope/University of Chicago.
- Post, J.M. (2004). *Leaders and Their Followers in a Dangerous World*. Cornell University Press.
- Scheidel, W., & Friesen, S.J. (2009). The size of the economy and the distribution of income in the Roman Empire. *Journal of Roman Studies*, 99, 61–91.
- Strauss, B. (2015). *The Death of Caesar: The Story of History's Most Famous Assassination*. Simon & Schuster.
- Suetonius. (c. AD 121). *The Twelve Caesars: Divus Julius*. Trans. Robert Graves. Penguin Classics.
- Wang, X., & Guo, Y. (2019). Leadership trait analysis of Chinese Communist leaders. *Journal of Chinese Political Science*, 24(3).
-

Editor's note

This report was researched and written with the assistance of AI tools for structural analysis, data synthesis, and drafting. All citations were verified against primary and peer-reviewed secondary sources. Leadership Trait Analysis scores were coded by the author using the Hermann (1999) rubric applied to primary source texts; they represent expert coding, not computational output. Roman GDP estimates are drawn directly from published academic sources and presented with their documented uncertainty ranges. The Psychonomics framework —

applying behavioral economics methodology to historical and contemporary leadership profiles — is an original contribution of Artometrics. Readers are encouraged to follow the source citations and form their own conclusions.

— Artometrics Editorial

[View source on GitHub](#)

Beyoncé and the Cost of Controlling the House

IN THIS REPORT

- [Research question](#)
- [Fast facts](#)
- [Data and method](#)
- [Leadership Trait Analysis](#)
- [The economic fingerprint](#)
- [The control network](#)
- [The affiliation machine](#)
- [Spectacle density](#)
- [Limitations](#)
- [Conclusion](#)
- [References](#)
- [Editor's note](#)

Economics comes from the Greek *oikonomia* — the management of the house. Psychonomics asks the prior question: who is managing the soul of the house, and what does their inner architecture cost everyone living inside it? The founding Artometrics case was Julius Caesar: a leader who treated the self as the institution, extended trust freely, and was removed by the highest-centrality nodes in his own patronage graph.

Beyoncé Knowles-Carter is the entertainment-industry case for the same method. Born in 1981, solo debut album year 2003 in the Artometrics Rolling Stone file, she is not a head of state. She is a head of catalog — and of Parkwood Entertainment, the firm that makes authorship itself the load-bearing wall. The psychonomic claim of this report is specific: where Caesar's theory was *magnanimous personal power*

without redundancy, Beyoncé’s theory is *owned control of the creative house*, paired with an affiliation-dense public voice and a spectacle-heavy measurable catalog.

Research question

If we score Beyoncé with the same Leadership Trait Analysis scaffold used for Caesar, and place those scores beside album-canon economics, a control network, and the Artometrics lyric extract, what theory of power does the profile encode — and what does it cost?

Fast facts

The numbers that set the scale for this report:

0.91 Task Orientation (T) — expert-coded LTA high point

0.88 Belief in Ability to Control Events (B)

#32 Lemonade on Rolling Stone 500 (2020 list)

186 Billboard weeks — 2013 self-titled album

8,000 Lyric lines in Artometrics Beyoncé extract

543 vs 5 Tokens: “love” vs “power” in that extract

Data and method

This report applies Margaret Hermann’s Leadership Trait Analysis (LTA) framework — seven variables scored 0–1.0 against a world-leader reference baseline: Belief in ability to control events (B), Conceptual complexity (CC), Distrust of others (D), In-group bias (IB), Need for power (P), Need for affiliation (A), and Task orientation (T). The Caesar Psychonomics report used the same scaffold. Comparative baseline values match that file’s Hermann reference average.

Beyoncé’s LTA scores are expert-coded for this series from: (1) public interviews and acceptance speeches; (2) documented firm practices around Parkwood and surprise-album strategy; (3) Homecoming-era work-process evidence in the Artometrics Netflix/Emmy rows; and (4) lyric-performance proxies from the 8,000-line extract. They are *not* automated NLP scores and *not* intelligence-community validated assessments. They are Psychonomics profiles: transparent expert coding with stated bias risk.

Observed Artometrics tables: lyric extract (‘beyonce-taylor-lyrics’), Rolling Stone albums file (Lemonade; self-titled), Emmy awards rows and Netflix title row for *HOMEcomings: A film by Beyoncé*. **Derived for this report:** LTA coding table, qualitative control network, lexicon and catalog-composition charts. Supporting CSVs ship under /data/articles/beyonce-the-psychonomics-of-control/.

Leadership Trait Analysis

High control and task, mid-high distrust, elevated affiliation — not Caesar’s shape

Beyoncé’s coded profile spikes on **Task Orientation (T = 0.91)** and **Belief in Ability to Control Events (B = 0.88)**, with **Need for Power (P = 0.86)** and **Conceptual Complexity (CC = 0.82)** also well above the world-leader baseline. **In-group bias (IB = 0.78)** and **Need for Affiliation (A = 0.74)** sit high. **Distrust of others (D = 0.62)** is above baseline (0.48) — the opposite direction from Caesar’s D = 0.31.

That asymmetry is the psychonomic argument. Caesar’s dangerous triangle was ceiling power + floor distrust: convert everyone, absorb the Republic into the self, die by the inner ring. Beyoncé’s triangle is ceiling task/control + mid-high distrust + high affiliation: speak in “we,” build a fandom institution, but do not outsource authorship governance to rivals you have pardoned into the center. The public voice is warm; the firm boundary is cold.

Power here is not primarily the language of the lyric tokens — “power” appears only five times in the extract — it is the right to decide what the house releases, when, and under whose name. Surprise drops and visual-album framing are behavioral B and T: the schedule is an instrument of control, not a concession to the industry calendar.

The economic fingerprint

Self-titled wins time-on-chart; Lemonade wins the canon list

The Artometrics Rolling Stone albums file carries two Beyoncé studio albums on the 2020 RS500 list. The 2013 self-titled album: Billboard peak #1, 186 weeks on chart, Spotify popularity 75, RS rank #81, artist age-at-list ~32. *Lemonade* (2016): peak #1, 87 weeks, Spotify popularity 73, RS rank #32, age-at-list ~35. Debut album year in that file: 2003.

This is the entertainment analogue of Caesar's GDP paradox, inverted in time scale. Commercial endurance (self-titled's years on chart) and canon elevation (*Lemonade*'s leap up the RS list) are not the same dependent variable. Psychonomics cares about both: one measures how long the market held the object; the other measures how the critical house re-ranked authorship after the visual-album claim. *Lemonade* is the album where the soul of the house became hardest to separate from the political and personal narrative it staged — and the canon file rewarded that fusion.

Later reinvention eras (Renaissance, Cowboy Carter) matter for the fame-path thesis in the MusicBrainz report — artists who become eras rather than moments — but they are not yet rows in the RS500 extract. They are timeline context, labeled as such.

The control network

Parkwood is the institutionalized self

Caesar's patronage chart had every edge terminating on a person. Beyoncé's chart has a second center of gravity: **Parkwood Entertainment** — the firm that converts personal authorship into an institution that can sign, stage, and film without surrendering the name. Family (JAY-Z, Blue Ivy), formative group (Destiny's Child circle), major-label distribution, live spectacle (Coachella → Homecoming), feature peers, and the Beyhive all connect — but governance weight sits on the Parkwood edge.

Artometrics' Emmy file shows the auteur extension: *Homecoming: A Film By Beyoncé* carries nominee credit rows for Beyoncé Knowles-Carter across directing, writing, music direction, and executive producer / performer categories (2019), with Parkwood and Netflix in the distributor/producer fields. That is not "celebrity appears in credits." That is the network claiming multiple authorship roles at once — B and T made bureaucratic.

The failure mode, if there is one, is different from Caesar's. It is not three pardoned senators with knives. It is concentration risk: when the firm and the person are coextensive, reputational or capacity shocks hit the whole house. Mid-high distrust ($D = 0.62$) is the wall that keeps rival nodes out of that core. High affiliation ($A = 0.74$) is how the outer ring stays emotionally enrolled while remaining structurally peripheral to governance.

The affiliation machine

The lyric extract speaks “love” and “baby,” not “power”

The packaged Artometrics extract holds **8,000** lyric lines across **139** song entities (~61.8k tokens, ~3,150 types). Top content signals are relational and imperative: **love 543, baby 457, boy 283, girl 239, night 104**. Identity/place markers appear at lower but non-trivial rates (**black 59, texas 18**). Abstract power language is scarce (**power 5, queen 9**).

This is where LTA's Need for Affiliation becomes audible without becoming the governance story. The songs recruit; the firm rules. Repeated lines in the extract are hook/choreography loops — refrain engineering — consistent with task orientation applied to bodies in a room, not with treatise-like power vocabulary. Psychonomics reads that split as intentional architecture, not contradiction.

Spectacle density

Features, live cuts, and remixes thicken the measurable catalog

Among 139 song entities in the extract: **35** carry Ft. credits, **35** are live/awards/“Experience”-type performance documents, **19** are remixes, and **64** fall outside those flags. Top line-count entities include VMA performance text and a Destiny's Child medley — long spectacle documents, not three-minute singles. Collaboration titles skew toward JAY-Z, with additional peers (Kanye West, Pharrell Williams, Nicki Minaj, and others) as catalog edges rather than co-governors.

For Psychonomics, spectacle density is evidence of T and B under live conditions: the house is managed by someone who treats the stage as a controllable system (Homecoming's whole point) and treats the catalog as expandable through remix and feature economies without surrendering the center.

Limitations

LTA was built for political leaders with large spontaneous-speech corpora. Extending it to a living entertainer requires caveats: public interviews are media-trained; lyrics are aesthetic objects; artist statements are strategic. Scores may overstate Task Orientation (work ethic is part of the brand) and understate private distrust. They should be read as the signature projected and institutionally enacted — which is what Psychonomics asks for — not as a full inner life.

The lyric file is an extract, not a complete discography. The published comparison article elsewhere on the site cites a larger merged Beyoncé-Taylor table; the downloadable Artometrics CSV used here is the 8,000-row Beyoncé-only pack. Rolling Stone rows cover two albums on one list year. Emmy/Netflix evidence is Homecoming-centered. The control network is qualitative, in the same methodological family as the Caesar patronage chart — documented roles, not a scraped graph. Renaissance and Cowboy Carter are narrative era markers pending richer quantitative tables.

Conclusion

Beyoncé's Psychonomics profile is a high-control, high-task, affiliation-forward leadership shape with Parkwood as the institutionalized self. The economic fingerprint splits commercial longevity (self-titled's 186 Billboard weeks) from canon elevation (Lemonade at RS #32). The lyric machine sells belonging; the firm keeps the keys. Compared with Caesar, the innovation is structural: distrust is high enough to keep knives out of the inner ring, and the house is built to be owned, not merely personally held.

Economics can track streams and chart weeks. Psychonomics asks who is managing the soul of the house — and what it costs. For Beyoncé, the cost of control is vigilance and concentration; the return is authorship that still has her name on the door when the era changes.

References

- Hermann, M.G. (1999). *Assessing Leadership Style: A Trait Analysis*. Social Science Automation.
- Artometrics. (2026). *Julius Caesar and the Cost of Managing Power* — Psychonomics method template. artometrics.com/caesar-the-psychonomics-of-emperor-julius
- Artometrics. (2026). *Beyoncé and Taylor Swift, Line by Line* — lyric extract. artometrics.com/beyonce-taylor-lyrics
- Artometrics Rolling Stone albums dataset — Lemonade (2016) and Beyoncé (2013) rows.
- Artometrics Emmy awards + Netflix titles datasets — *HOMEcomings: A film by Beyoncé*.
- Artometrics. (2026). *MusicBrainz pop fame mechanics* — era-maker framing. artometrics.com/musicbrainz-pop-fame-mechanics
- Data Science Learning Community. (2020). TidyTuesday Beyoncé & Taylor Swift lyrics release (source lineage for lyric extract).
- Supporting tables: /data/articles/beyonce-the-psychonomics-of-control/ (beyonce_lta.csv, beyonce_eras.csv, beyonce_network.csv, beyonce_lyric_summary.csv).

Editor's note

This is the second full Psychonomics profile in the Artometrics series (after Caesar). LTA scores were coded by the editorial desk using the Hermann rubric; they are expert judgments with disclosed corpora, not model outputs presented as measurements. AI tools assisted structuring and drafting under human review. Readers should treat the radar as an argument with numbers — and follow the limitations before treating any single trait as destiny.

How Long Do CEOs Last Before They Leave?

CEO tenure is usually narrated as drama: a firing, a founder exit, a sudden interim. The TidyTuesday departures extract used here holds **9,423** records with a median max tenure (ceodb) of **1.00** and a high of **4.00**. Interim is the most common co-CEO/interim label in the file.

That median of 1.00 is the calibration point: in this coding, many departure spells are short. The charts ask which company names show longer max tenures, how interim versus co-CEO labels differ, and whether tenure fields move together.

Fast facts

The numbers that set the scale for this report:

9,423Records in the working dataset

1.00Median Max tenure ceodb

4.00Highest observed Max tenure ceodb

PHOTRONICS INCTop Coname by Max tenure ceodb

InterimMost common Interim coceo

Data and method

The source is the TidyTuesday release from 2021-04-27 (departures.csv). After cleaning, 9,423 rows remain.

Max tenure ceodb is the primary ranked metric. Interim/co-CEO spelling variants appear as separate categories in places and should be read as label noise as well as structure. Charts are Plotly JSON with PNG fallbacks.

Longer Max Tenures Are Rare in the Company Ranking

Max tenure ceodb by Coname

Stewart Information Services and Conversant Inc lead the company view at **2.50** on max tenure ceodb. A wider set of firms — Avatex, Intrepid Potash, Bergen Brunswig, SEACOR, Micro Warehouse, RH — cluster at **2.00**.

Relative to the file median of 1.00, even these “leaders” are only modestly longer spells. Extended CEO tenure is the exception in this extract, not the rule.

The Top Dozen Caps Out Quickly

CONVERSANT INC leads at 2.50 — 2.00 marks the median among the top dozen

Conversant Inc appears among the leaders at **2.50**, and the median among the top dozen is **2.00**. There is no long tail of decade-scale max tenures in this particular ranking metric.

Cite 2.00 as the elite-club median for max tenure ceodb here — double the overall median, still short by popular mythologies of imperial CEOs.

Interim Labels Dominate the Category Mix

Max tenure ceodb by Interim coceo

Interim accounts for **218** boxed observations with a median max tenure of **1.0** and a mean near **1.23**. CO-CEO shows **85** rows with the same median of 1.0 but a slightly higher mean (~**1.4**). Smaller capitalization variants of co-CEO are tiny samples.

The archive is thick with interim spells. That is a succession-system fact: departures often pass through temporary authority before a permanent appointment.

Gaps to the Median Are Mostly Flat

Max tenure ceodb vs median by Interim coceo

Most interim/co-CEO labels sit at a zero gap to the median on max tenure ceodb in this chart, with only a thinly populated CO-ceo variant showing a positive gap of **1.0**.

Category labels do not produce a dramatic tenure hierarchy here. The bigger story remains how many spells are coded interim at all.

Tenure Fields Track Closely Within Labels

Max tenure ceodb vs Tenure no ceodb

Max tenure ceodb versus tenure no ceodb, colored by interim/co-CEO label, largely follows a tight correspondence — points hug aligned values, especially in the interim cloud (n=218).

Where the two tenure fields disagree, coding definitions matter more than boardroom legend. The scatter is a consistency check on the file's tenure construction.

What this file cannot tell you

Tenure fields are database constructions, not calendar biographies. Spelling variants of co-CEO fragment categories. Firm name changes and dual-class structures can split identities.

A median of 1.00 may reflect coding granularity as much as “CEOs only last a year.” Read the metric definition before converting it into a headline about corporate instability.

What to take away

In 9,423 departure records, max tenure ceodb centers at 1.00, with company-level leaders only reaching 2.0–2.5. Interim is the dominant succession label.

The citable story is institutional: departure files are full of short and interim spells, not a museum of decades-long reigns.

Sources

Data Science Learning Community. (2021). *TidyTuesday: CEO Departures*. <https://raw.githubusercontent.com/rfordatascience/tidyuesday/main/data/2021/2021-04-27/departures.csv>

How Concentrated Is Wealth at the Top of the Distribution?

Income distribution is how a society counts who sits where. This file holds 2,916 records spanning 1967–2019, with a median income-distribution value of 10.9 and a high of 27.2. Black Alone appears in the fact-box race ranking; Under \$15,000 is the most common income bracket label.

The charts track how the median moved, which race categories lead, how brackets sit relative to the median, and how income distribution relates to income median. The calibration point is 10.9 — the center of the distribution field in this extract.

Fast facts

The numbers that set the scale for this report:

2,916Records in the working dataset
10.9Median Income distribution
27.2Highest observed Income distribution
Black AloneTop Race by Income distribution
1967–2019Year span covered in the file
Under \$15,000Most common Income bracket

Data and method

The source is the TidyTuesday release from 2021-02-09 (R for Data Science community). The working file contains 2,916 rows and 10 columns after merging available tables in the week folder. Income distribution is the primary metric; race and income bracket are categorical axes; income median appears in the scatter.

Medians are used for robustness across skewed economic series. Index-style fields are excluded from metric selection.

The median income-distribution marker drifted upward

Median Income distribution Over Time

Median income distribution rises from 10.9 in the opening period to 11.2 at the close — a small but measurable lift around the file median. The center of this field did not explode; it edged higher across the long window.

Modest median drift can coexist with large bracket-level gaps. The trend chart reports the center; the gap chart reports the structure underneath.

Hispanic (Any Race) leads the charted race ladder

Hispanic (Any Race) leads at 12.0 — 10.7 marks the median among the top dozen

Hispanic (Any Race) leads at 12.0, while 10.7 marks the median among the top dozen race categories in this cut. Black Alone remains a fact-box landmark for the ranking rule used there; the leaders chart makes the numeric top of this aggregation explicit.

The ladder is tight near the center of the file — leaders sit close to the overall median of 10.9 rather than in a different statistical universe.

Income brackets carve different distribution bands

Income distribution by Income bracket

Category boxes by income bracket show whether the income-distribution field is shared or contested across dollar bands. Under \$15,000 is the most common bracket label by frequency; the boxes ask how the metric spreads inside each band.

Bracket structure is the skeleton of inequality analysis in this file. A single median of 10.9 hides those band-level differences.

Mid brackets clear the median; high brackets trail in this cut

Income distribution vs median by Income bracket

\$50,000 to \$74,999 sits 6.60 above the median; \$150,000 to \$199,999 trails by 6.70. Those signed gaps are among the starkest bracket contrasts in the extract.

Trailing or leading the median here is a statement about the income-distribution field's values by bracket, not a casual claim about who is "richer" in narrative terms. Cite the signed distances as structural off-sets.

Distribution values and income medians form joint clusters

Income distribution vs Income median

Plotting income distribution against income median shows clusters that averages erase. Race-year points co-locate in patterns that a single summary cannot hold.

The scatter is relational: it shows how the two income fields move together across the 1967–2019 extract.

What this file cannot tell you

Community-cleaned TidyTuesday snapshots are not live Census microdata APIs. Missing values, race and bracket labeling choices, and 1967–2019 coverage limits apply. Merged tables may fan out or duplicate rows when join keys are imperfect.

Findings describe this wealth-and-income extract — structural signals about the income-distribution field — not a complete wealth account including assets and debt.

What to take away

Income distribution in this file centers on 10.9, edges up to 11.2 at the median, and still shows sharp bracket gaps — mid bands above the median, top dollar bands trailing in this metric's signed distances.

Race ladders and the distribution–median scatter add group and joint structure. Together they map concentration and composition without reducing fifty years of income data to a single slogan.

Sources

Data Science Learning Community. (2021). *TidyTuesday: Wealth and Income*. https://raw.githubusercontent.com/rfordatascience/tidyuesday/main/data/2021/2021-02-09/income_distribution.csv

Editor's note

Artometrics data report from the TidyTuesday research pipeline. Charts and aggregates are reproducible from the embedded exhibits and public source files.

[View TidyTuesday source on GitHub](#)

When Louis Vuitton Beat Molly Tea in Court — and Lost the Chinese Internet

IN THIS REPORT

- [Research question](#)
- [Fast facts](#)
- [Data and method](#)
- [What the court actually ordered](#)
- [A larger litigation machine](#)
- [Resemblance is not provenance](#)
- [Secondary-market signal, not a balance sheet](#)

- Two frames, same facts
- Limitations
- Conclusion
- References
- Editor's note

Chinese short-form video turned a Suzhou trademark judgment into a morality play: a French luxury house “stole” a Chinese pattern, trademarked it, then sued a domestic tea chain. The court record is narrower. Louis Vuitton Malletier sued Shenzhen Molly Tea Catering Management Co., Ltd. and a franchise outlet over a four-petal floral device used as brand identity. The first-instance court found infringement of seven registered marks and awarded ¥10.3 million. Molly Tea says it will appeal.

This report separates three layers that the viral clips fuse: (1) what the judgment and CNIPA history show; (2) how the case sits inside Louis Vuitton’s wider China enforcement pattern; and (3) why a legally ordinary win became a cultural-economy event — including what resale listings can and cannot prove.

Research question

When a foreign luxury registrant wins a China trademark case against a popular domestic consumer brand, what do the filings establish — and what does the public backlash measure instead?

The observational question is whether the Molly Tea dispute is best read as trademark doctrine (similarity, association risk, prior CNIPA refusal) or as a heritage-and-optics crisis (motif commons, nationalist brand politics, soft luxury demand). Those readings share images; they do not share a dependent variable.

Fast facts

The numbers that set the scale for this report:

¥10.3M First-instance award (damages + costs)
 7LV flower trademarks held infringed
 ~1,691LV-linked China TM cases ~5 years (media-cited)
 56Reported new LV China filings, H1 2026
 6Platforms ordered for corrective statements
 AppealJudgment is first instance, not final

Data and method

This is a documentary case study, not a full docket scrape. Primary legal facts are taken from Chinese and English IP reporting on case **(2025) 苏05民初617号** (Suzhou Intermediate People's Court), including damages figures circulated from a leaked judgment page, plus AP and specialist IP write-ups. Enforcement volume figures (~1,691 records / 56 H1 2026 filings) are *media-cited company-search tallies* (Qichacha via China Briefing / Global Times), not a cleaned Artometrics microdata file — treat them as order-of-magnitude context.

Resale examples (Carryall, Neverfull, Onthego) are *illustrative listings* reported by Chinese secondary-market coverage after the ruling. They are not a panel of matched pre/post prices. Where we show an “editorial index” comparing court-frame vs public-frame weights, that chart is an interpretive device, not observed survey data.

Observed vs derived: award amounts, trademark counts, case number, and platform apology list are observed from reporting on the judgment. “Interpretive weight” scores and any ratio commentary on resale discounts are derived editorial indices. Viral social claims (“LV = lavatory,” “carrying LV = China hater”) are treated as discourse, not prevalence estimates.

What the court actually ordered

The award is ¥10.3 million — and it is not final

Reporting on the judgment describes a first-instance order for Shenzhen Molly Tea to cease infringing use of the disputed four-petal device, pay **¥10 million** in economic damages and **¥300,000** in reasonable enforcement costs, and publish corrective statements on six surfaces: official website, Weibo, WeChat Official Account, WeChat Mini Program, Xiaohongshu, and Douyin. A franchise defendant was described as jointly liable up to ¥100,000. Molly Tea has publicly said it will appeal.

Two procedural facts matter more than the meme version. First, Molly Tea had already tried to register a four-petal logo; CNIPA refused application no. **81049590** (Class 43), citing prior LV flower marks, and the refusal review became final. Second, the court's association logic sits in a market where luxury-FMCG collaborations are common — so

“nobody confuses a handbag with milk tea” is not the only test. The legal question was whether the logo could suggest sponsorship or commercial linkage.

A larger litigation machine

One tea chain sits inside nearly 1,700 China trademark records

Chinese coverage places Molly Tea inside a multi-year Louis Vuitton enforcement campaign: on the order of **1,691** trademark-infringement records over about five years, with **56** new filings reported in the first half of 2026. Recoveries across many matters are described as modest relative to global luxury revenue — deterrence and anti-dilution, not a profit center.

That context cuts both ways. It undercuts the idea that Molly Tea was a one-off “bullying” novelty: LV litigates flower-device conflicts as a portfolio habit. It also explains why the public read this case as the tip of a pattern — especially when parallel administrative fights with CNIPA and other small defendants entered the same news cycle. Comparative China brand fights (for example, long-running New Balance mark disputes, or fashion-heritage flare-ups such as the Dior / *mamianqun* controversy) show the same structural tension: first-to-file commercial exclusivity versus popular claims about shared cultural forms.

Molly Tea was founded in 2021, scaled a floral visual system through stores and packaging, hit a CNIPA wall in 2024, was sued in May 2025, and lost at first instance in late June 2026. The viral wave followed the judgment leak and apology-order optics — not the original complaint filing.

Resemblance is not provenance

Tang baoxiang comparisons are real; they are not a trademark defense by themselves

Xiaohongshu and state-media posts placed LV’s four-petal monogram flower beside Tang-dynasty decoration — including rosewood *pipa* inlay and quatrefoil window lattice in classical gardens such as Suzhou’s Humble Administrator’s Garden. Historians describe *baoxiang* motifs

as composite ornamental vocabularies drawn from lotus, peony, and related flora. Persimmon-calyx geometry is another folk visual rhyme circulating in the same threads.

Those comparisons establish **visual kinship in a shared decorative commons**. They do not, by themselves, cancel a registered commercial graphic under China's Trademark Law. LV's own origin story for the 1896 Monogram, as relayed in LVMH materials and wire reporting, cites neo-gothic ornament and Japonism — a nineteenth-century European fashion for East Asian decorative influence — not a court finding that the mark is a Tang-era replica.

The honest formulation: modern IP systems let firms fence *specific* commercial embodiments of shapes that also exist, in looser form, in heritage arts. The Chinese internet's objection is that fencing feels like enclosure of something older than the registrant. The court's objection to Molly Tea is that the chain's standardized black-and-white four-petal mark tracked LV's protected device too closely after CNIPA had already said so.

Secondary-market signal, not a balance sheet

Resale listings moved; attribution is messy

Chinese secondary-market coverage after the judgment described quieter boutiques and softer ask prices — for example, a Carryall small near ¥19,000 new on Dewu against ~¥23,000 retail, and deeper discounts on used Neverfull and Onthego units. Dealers also described a surge of sellers. That is a real market *mood*.

It is not yet a clean causal estimate. Louis Vuitton's China demand was already softening in the broader luxury slowdown; dealers can talk down bids when headlines are bad; listing prices are not cleared transactions. The viral claim that bags are “down 10–20% because of Molly Tea” overstates what public listing anecdotes can identify. Treat post-ruling resale chatter as a coincident signal inside a pre-existing demand slide.

Two frames, same facts

Court doctrine and internet dignity are different dependent variables

In the court frame, registered rights, CNIPA prior refusal, and association risk dominate. In the public frame, Tang resemblance, nationalist brand optics, and “foreign giant vs domestic chain” narrative dominate. East Is Read’s media monitoring put related posts in the hundreds of thousands with Weibo topic pages in the hundreds of millions of views — order-of-magnitude evidence of attention, not of legal error.

So: is the hate “fair”? Fairness is the wrong metric if it forces one answer. Under trademark doctrine, LV asked a Chinese court to do what Chinese first-to-file practice is designed to do — and got a first-instance win. Under brand strategy in a prestige market that runs on story and belonging, LV converted a winnable IP dispute into a highly shareable parable about cultural enclosure. Both can be true at once.

Limitations

We do not have a full public official reporter text for every page of the judgment — only consistent secondary reporting and a circulated damages page. Enforcement volume figures are media-cited search tallies, not a verified Artometrics dataset. Resale examples are illustrative listings, not a difference-in-differences study. The “two frames” chart is an editorial index. Social claims about bathroom memes and “China hater” labeling are discourse samples, not survey prevalence.

Conclusion

Molly Tea is a first-instance trademark defendant ordered to pay ¥10.3 million and apologize on six platforms; the case is on a path to appeal. Louis Vuitton’s legal theory sits inside a large China enforcement machine. The internet’s theory sits inside a longer argument about who may exclusively commercialize motifs that also live in Chinese decorative history. The Shorts that spread the story are useful as a map of popular feeling. They are a poor substitute for the docket.

If there is a citable lesson for creative-industry readers, it is this: in China as elsewhere, winning the register and winning the room are different sports. Luxury brands that litigate flower geometry against domestic consumer heroes should price reputation the way they price leather — as inventory that can mark down overnight.

References

- AP News — “Louis Vuitton court victory against Chinese tea chain stirs up a copyright debate” (2026). apnews.com
- China IP Law Update — “Louis Vuitton Scores 10 Million RMB in a Controversial Trademark Infringement Dispute Against Molly Tea” (July 2026). chinaiplawupdate.com
- D’Andrea & Partners — “Louis Vuitton v. Molly Tea: Chinese Court Rules on Trademark Infringement in the Beverage Sector.” dandreapartners.com
- China Briefing — “Trademark Protection in China: Legal Victory Is Only Part of the Strategy.” china-briefing.com
- Global Times — “LV vs Molly Tea case exposes modern IP’s cultural blind spot.” globaltimes.cn
- Jiemian / Xinhua / Southern Metropolis secondary reporting on judgment date, appeal statements, and damages.
- Artometrics ethics — [evidence tiers](#).

Editor’s note

This report was prompted by a widely shared YouTube Short from [@urchineseuncle](#) (Shushu) framing the Louis Vuitton–Molly Tea dispute for an English-speaking audience. Artometrics’ contribution is complementary: court identifiers, CNIPA refusal history, enforcement-volume context, motif-vs-trademark distinctions, and limits on resale anecdotes. Credit to the original explainer for surfacing the story; the analysis and any errors here are ours.

AI tools assisted research organization and drafting under human editorial review. Claims were checked against named secondary legal and wire sources listed above.

How U.S. Incarceration Rates Diverged by State

U.S. incarceration is often narrated as a single national curve. County- and state-level population files show something more fractured: regional regimes that diverge for decades even while they share the same federal criminal-justice vocabulary.

The TidyTuesday working extract used here holds **100,000** records spanning **1970–2016**. Median population in the analytic columns sits at **3,553**; the highest observed population value exceeds **6.8 million**. California leads the state rollup by population, and the South is the most common region label in the file.

Population is not identical to incarceration rate — but in this release it is the spine that lets pretrial counts, regional gaps, and long-run change be compared without pretending every county is the same size.

Research question

What does the county-and-state structure of the U.S. incarceration file reveal that a single national mass-incarceration curve would hide? This report asks how population, region, and pretrial population sit together in the 1970–2016 extract, and where the file points toward divergent local regimes rather than one uniform national process.

The question is observational: which geographies dominate the row structure, how far regional medians sit from the file center, and whether pretrial burden scales cleanly with population. It does not treat population as an incarceration rate; it treats it as the denominator context necessary before any rate-based claim can be trusted.

Fast facts

The numbers that set the scale for this report:

100,000Records in the working dataset
3,553Median Population
6,801,437Highest observed Population
CATop State by Population
1970–2016Year span covered in the file

SouthMost common Region

Data and method

The source is the TidyTuesday release from 2019-01-22 (R for Data Science community). The working file contains 100,000 rows and 10 columns after assembly — state, region, year, population, and pretrial population among the fields carried forward.

Medians stabilize skew from giant counties and city systems. Charts export as Plotly JSON with PNG fallbacks. Because the table mixes geographies of very different scale, every ranking should be read with population context attached.

How the pattern changed over time

Median Population Over Time

Median population in the working series falls from about **4,123** in the opening period to about **3,359** at the close. That movement can reflect changing county coverage, urban-rural composition, and how incomplete early rows enter the sample — not a simple story that places got smaller.

The time chart is still essential. It forces the rest of the analysis to admit that the typical row in 1970 and the typical row in 2016 are not drawn from identical geographic mixes.

The 1970–2016 window covers the main period in which the United States moved from the pre-1970 jail-and-prison baseline into what scholars such as the National Research Council and Bureau of Justice Statistics describe as the era of mass incarceration. It includes the years after the 1971 declaration of the federal war on drugs, the sentencing-policy changes of the 1980s, and the long plateau and partial retrenchment after the 2000s.

Because this file is row-based, the median population line is partly a coverage diagnostic. A county with a small jail, a large urban system, and a state-level aggregate can coexist in the same extract if the source tables carry them forward. The trend line therefore tells readers how the analytic unit changes over time before it tells them anything about punishment intensity.

Who sits at the top

DC leads at 190,209 — 24,944 marks the median among the top dozen

On the leader cut highlighted in the file, **DC** leads at **190,209**, with **24,944** as the median among the top dozen. Dense jurisdictions and large states occupy the same head for different reasons — capital-city concentration versus sheer statewide population.

California's fact-box lead on state population and DC's lead on this chart can both be true. They answer different aggregation questions. Confusing them produces false arguments about which place is most incarcerated rather than which place is largest or densest in the extract.

The District of Columbia is a useful warning case because it is not a state, not an ordinary county, and not a simple municipal equivalent. Its criminal-justice system has long involved federal oversight, local jail custody, and post-sentencing arrangements that differ from state prison systems. A high DC entry in a population leader chart should trigger questions about geography and institutional definition before it triggers a national ranking claim.

California's presence in the state rollup points to a different scale issue. Los Angeles County, the California Department of Corrections and Rehabilitation, and statewide demographic size all shape the visibility of California in criminal-justice data. Large population does not automatically mean high rate; it means the denominator and administrative footprint are large enough to dominate many raw-count charts.

How regions spread

Population by Region

Regional box plots show whether population consensus is shared or contested across the South, Northeast, Midwest, and West. The South's frequency as the most common region label does not automatically mean it has the highest median population per row.

Spread within regions often exceeds spread between region medians. County inequality inside a region is part of the incarceration story, not noise around a regional average.

The regional labels matter because U.S. criminal justice is federal in vocabulary but local in operation. Sheriffs, county commissioners, state legislatures, prosecutors, bail rules, jail financing, and prison siting vary sharply across the South, Northeast, Midwest, and West. The South's high row frequency in the file is consistent with its large county count and centrality in incarceration scholarship, but it is not itself proof of a higher imprisonment rate.

Within-region spread is where the policy story begins. A rural Southern county, Harris County in Texas, Cook County in Illinois, and New York City's jail system do not share the same scale or pretrial process simply because they fit under regional labels. Box plots protect that variation from being erased by a single national average.

Who beats the median — and who trails

Population vs median by Region

On the gap chart, the **Northeast** sits **13,710** above the median; the **Midwest** trails by **1,276**. Those gaps capture the urban weight of Northeastern counties in the sample as much as any single policy variable.

Regional gaps are starting points for rate-based follow-ups, not final moral rankings. A high-population region can have lower incarceration rates, and a lower-population region can incarcerate more intensively per resident.

The Northeast's positive gap is plausibly shaped by the presence of dense metropolitan jurisdictions such as New York, Philadelphia, Boston, Newark, and surrounding counties, where raw population denominators are large. The Midwest's negative gap does not mean Midwestern incarceration is low; it means the population field in this working extract sits below the file median for that regional cut.

That distinction is not pedantry. Prison Policy Initiative, Vera Institute, and BJS reports routinely separate jail admissions, jail population, prison custody, and incarceration rates per 100,000 residents because raw headcounts can move in the opposite direction of rates. The gap chart is a routing map for those follow-ups, not a replacement for them.

Population and pretrial population

Population vs Pretrial population

Plotting population against pretrial population shows clusters that averages erase. Some large jurisdictions carry large pretrial counts in rough proportion; others deviate — the places where jail composition, not just city size, is doing extra work.

Pretrial population is one of the most policy-sensitive fields in modern jail debate. The scatter does not prove causation, but it flags where scale and pretrial burden travel together and where they separate.

Pretrial detention is legally different from sentenced imprisonment. People held pretrial have generally not been convicted in the case that keeps them in jail, which makes bail schedules, risk assessment tools, prosecutorial charging, and court backlogs central to the count. The Bureau of Justice Statistics has repeatedly documented that local jails hold a large share of unconvicted people, making this field a direct bridge between data and constitutional policy.

The scatter can therefore guide a more serious audit. Jurisdictions far above the population-pretrial trend are candidates for closer inspection of bail practices, case-processing time, and jail-capacity politics. Jurisdictions below the trend may have diversion programs, lower detention thresholds, different offense mixes, or simply different data coverage. The chart identifies candidates; it does not adjudicate causes.

What this file cannot tell you

Community-cleaned TidyTuesday snapshots are not live APIs. Missing values, spelling variants, and coverage limits apply. Population columns are not a substitute for official imprisonment rates per 100,000 residents.

Findings describe the file on hand — structural signals about how geographies and pretrial counts sit across 1970–2016 — not a complete causal history of mass incarceration.

What to take away

Incarceration trends in this extract are a story of divergent geographies sharing one national vocabulary. Median rows sit near 3,553 people; extremes exceed millions; regional gaps reach five figures above the median.

The citable caution is simple: size, region, and pretrial burden are related but not identical. Any claim about the U.S. incarceration rate that skips those splits is summarizing away the map.

Sources

Data Science Learning Community. (2019). *TidyTuesday: Incarceration Trends*. <https://github.com/rfordatascience/tidytuesday/tree/main/data/2019/2019-01-22>

Bureau of Justice Statistics. *Correctional Populations in the United States*. https://bjs.ojp.gov/library/publications/list?series_filter=Correctional%20Populations%20in%20the%20United%20States

National Research Council. (2014). *The Growth of Incarceration in the United States: Exploring Causes and Consequences*. <https://doi.org/10.17226/18613>

Vera Institute of Justice. *Incarceration Trends*. <https://trends.vera.org/>

Which States Show Up at the Ballot Box?

Turnout is democracy's volume knob. This file holds 936 election records spanning 1980–2014, with a median of 1,170,867 votes and a high of 132,609,063. United States appears as the top entity by votes in the fact boxes — a national aggregate sitting above the state ladder.

The charts separate trend, leaders, distribution shape, leader trajectories, and the relationship between votes and eligible voters. The calibration point is that 1.17 million median: below it, smaller states; above it, larger electorates and national totals.

Fast facts

The numbers that set the scale for this report:

936Records in the working dataset

1,170,867Median Votes

132,609,063Highest observed Votes

United StatesTop State by Votes

1980–2014Year span covered in the file

Data and method

The source is the TidyTuesday release from 2018-10-09 (R for Data Science community). The working file contains 936 rows and 7 columns after merging available tables in the week folder. Votes is the primary metric; eligible voters appears in the scatter; state (and national totals) are the entity axis.

Medians are essential because vote totals are right-skewed — mean 3,074,280 against median 1,170,867. Index-style fields are excluded from metric selection.

Median votes rose across three decades of elections

Median Votes Over Time

Median votes rises from 910,290 in the opening period to 1,387,622 at the close. Population growth, eligibility changes, and which elections enter the file all can lift the center; the observed trend in this extract is upward.

A rising median is the shared backdrop. Leader states and national totals amplify it; smaller jurisdictions still sit inside a shifting distribution.

The United States total dwarfs the state ladder

United States leads at 90,912,015 — 4,209,538 marks the median among the top dozen

United States leads at 90,912,015 votes in the leaders cut, while 4,209,538 marks the median among the top dozen. The national row is not comparable to a single state on equal footing; it is the aggregate ceiling in the file.

Fact boxes also list 132,609,063 as the highest observed votes — the extreme of the same national-scale phenomenon. State comparisons live far below that ceiling.

A right-skewed distribution with a heavy upper tail

Median 1,170,867 vs mean 3,074,280 — the shape is right-skewed

Median 1,170,867 versus mean 3,074,280 — the shape is right-skewed. Large states and national totals pull the average up. The top decile begins at 4,659,000; that is where defining high-volume cases live.

Skew is why turnout debates should lead with medians and rates, not raw means of vote totals alone.

Leading states reshuffle across election cycles

Top State Over Time

Leader trends show that leading names do not move in lockstep — some fade as others surge. Tracking votes over time separates sustained high-volume states from cycle-specific spikes.

Presidential versus midterm composition in the file can rearrange the ladder even when the long-run median is rising.

Votes scale with eligible voters — in clusters

Votes vs Eligible voters

Plotting votes against eligible voters shows clusters that averages erase. Larger eligibility pools produce larger vote totals, but the scatter's texture reveals how tightly that relationship holds across states and years.

The joint view is the right place to think about turnout as a rate problem without inventing a rate column the charts do not show: eligibility and votes move together, and the clusters show where.

What this file cannot tell you

Community-cleaned TidyTuesday snapshots are not live election boards. Missing values, jurisdiction naming variants, and 1980–2014 coverage limits apply. National and state rows may coexist in ways that require care when ranking.

Findings describe this voter-turnout extract — structural signals about vote totals — not a complete model of registration law, suppression, or mobilization.

What to take away

Vote totals in this file are right-skewed around a median of 1,170,867, rising from 910,290 to 1,387,622 at the median, with national rows reaching into the tens and hundreds of millions.

State ladders reshuffle over time, and votes move with eligible voters in clusters — a volume atlas of American elections from 1980 through 2014.

Sources

Data Science Learning Community. (2018). *TidyTuesday: US Voter Turnout*. https://raw.githubusercontent.com/rfordatascience/tidyuesday/main/data/2018/2018-10-09/voter_turnout.csv

Editor's note

Artometrics data report from the TidyTuesday research pipeline. Charts and aggregates are reproducible from the embedded exhibits and public source files.

[View TidyTuesday source on GitHub](#)

Who Ruled Rome Longest—and Did Dynasty Predict Tenure?

Sixty-eight emperors sit in this reign-length file — enough to ask a hard question about imperial power: did belonging to a dynasty buy more time on the throne, or did longevity come from something messier? The calibration point is a median reign of 6.50 years. That is not a long reign. It is the mathematical center of a political system that chewed through rulers faster than most modern institutions replace CEOs.

The ceiling in the file is Constantine the Great at 31.0 reign years. Gordian is the dynasty label that appears most often. Between those two facts lies the real story: a few men held power for decades, while entire dynastic brands recycled short, fragile tenures. The charts that follow separate name-level longevity from dynasty-level structure so the two are not confused.

Research question

Did dynastic identity make Roman emperors more durable, or did individual political survival overwhelm house-level labels? This report asks whether reign length concentrates around a few exceptional rulers, whether dynasties show distinct tenure distributions, and how much aggregate imperial time is carried by the longest-reigning names.

The question is observational because this 68-record TidyTuesday file is not a full constitutional model of Rome. It uses recorded reign spans to compare names and dynasties, then treats the median emperor's 6.50-year tenure as the baseline against which Constantine, Nerva-Antonine, and Gordian entries can be read.

Fast facts

The numbers that set the scale for this report:

68Records in the working dataset

6.50Median Reign years

31.0Highest observed Reign years

Constantine the GreatTop Name by Reign years

GordianMost common Dynasty

Data and method

Reign spans are derived from start and end dates in the emperors table. Negative or zero spans are excluded, so the working metric is observed tenure in years, not a modeled life expectancy. The source is the TidyTuesday Roman Emperors release (2019-08-13), a community-cleaned historical table rather than a live prosopography API.

Medians are preferred over means because reign distributions skew — a handful of long rulers pull averages up while the typical emperor sits near that 6.50-year center. Index-style fields (row numbers, sequential IDs) are kept out of metric selection so charts answer questions about tenure, not about how the spreadsheet was ordered.

Constantine's thirty-one years set the name-level ceiling

Reign years by Name

Among named emperors in the breakdown, Constantine the Great leads at 31.0 reign years. At the other end of the plotted leaders, Valentinian II anchors the low end of that top group at 17.0 — still far above the file median of 6.50, which is the point: the chart of leaders is a curated upper slice, not the full distribution.

Grouping by name makes the catalog's major entities comparable on a single scale. The question is not who is famous, but who lasted. The gap between Constantine and the rest of the top tier already hints that longevity was exceptional rather than routine.

Constantine's 306–337 CE reign sits at a hinge in Roman institutional history: civil war, Tetrarchic succession, Christian legalization after the Edict of Milan, the Council of Nicaea in 325, and the foundation of Constantinople all fall inside or around his rule. The dataset reduces that complexity to 31.0 years, but the named entity carries enough historical context to explain why the ceiling is not just a spreadsheet oddity.

The top dozen still sits far above the median emperor

Constantine the Great leads at 31.0 — 20.0 marks the median among the top dozen

Constantine the Great leads the ranked leaders chart at 31.0. Among that top dozen, the median is 20.0 — three times the full-file median of 6.50. Even the “typical” elite tenure in this slice is a different political animal from the typical emperor in the dataset.

That separation matters for interpretation. Leaderboards describe the right tail. They do not describe Rome's everyday succession reality, which lived closer to half a decade than to two decades.

The 20.0-year median among the top dozen contains the familiar names that make imperial history feel more stable than it was: Augustus, Tiberius, Hadrian, Antoninus Pius, Marcus Aurelius, Constantine, and other long incumbents dominate classroom memory. The median of the full file pulls the reader back toward usurpers, assassinations, co-rule, civil wars, and short-lived crisis reigns.

Dynasty boxes show contested, not uniform, tenure

Reign years by Dynasty

When reign years are split by dynasty, the boxes do not collapse into a single shared consensus. Some houses stretch higher; others sit tighter and lower. The chart's job is to show whether dynasty is a meaningful axis for tenure at all — and the spread says it is contested across tiers rather than deterministic.

Read the boxes as structural signals: belonging to a named dynasty did not automatically purchase Constantine-length stability. It bought a distribution, and those distributions differ.

The box plot therefore should not be read as a genetic theory of power. It is a way to ask whether succession packages mattered. The Nerva-Antonine adoption sequence, for example, has long been associated with high imperial stability, while the third-century crisis illustrates how military acclamation and rapid turnover can shred dynastic continuity.

Nerva–Antonine sits above the median; Gordian trails

Reign years vs median by Dynasty

Relative to the file median, Nerva-Antonine sits 12.5 years above; Gordian trails by 4.00. That is the cleanest dynasty-level contrast in the gap chart: one house systematically above the center of the distribution, another systematically below.

Gordian is also the most common dynasty label in the fact boxes — volume without longevity. Frequency of naming and length of rule are not the same political achievement.

The Nerva-Antonine result is historically intuitive because Nerva, Trajan, Hadrian, Antoninus Pius, Marcus Aurelius, and their succession network occupy the second-century high-water mark of imperial continuity in much Roman historiography. A positive gap does not prove the dynasty caused stability, but it aligns with the conventional account of fewer violent turnovers and longer administrative horizons.

The Gordian contrast points to the opposite problem. Gordian I and Gordian II ruled briefly in 238 CE during the Year of the Six Emperors, while Gordian III lasted longer but still belongs to a crisis environment shaped by military pressure and senatorial maneuvering. A repeated dynasty label can mark a famous rupture rather than a durable house.

Five names carry forty percent of aggregate reign years

Cumulative Reign years

The Pareto curve shows concentration: the top five name entries account for 40% of the aggregate reign years in the working file. A steep head means a small set of long rulers dominates the sum, while the long tail of shorter reigns fills out the catalog without moving the total as much.

That shape reframes the median. A 6.50-year center can coexist with a heavily concentrated aggregate if a few Constantine-scale tenures exist. Both facts are true; they answer different questions.

Concentration is a useful way to connect biography to institutions. Long reigns give rulers time to appoint governors, restructure armies, mint ideological programs, build capitals, codify religious settlements, and shape elite careers. Short reigns often leave the opposite signature: coinage, military proclamation, purge, assassination, or succession crisis.

What this file cannot tell you

This is a community-cleaned TidyTuesday snapshot, not a continuously updated classical-history database. Missing values, spelling variants, and week-of-export coverage limits apply. Reign years depend on how start and end dates were recorded and cleaned; disputed chronologies are not adjudicated here.

Dynasty labels are editorial categories in the source table. Findings describe the file on hand — structural signals about tenure and dynasty in this emperors extract — not a complete political history of imperial Rome.

What to take away

The file's central tension is simple: the median emperor lasted 6.50 years, Constantine the Great lasted 31.0, and dynasty membership did not flatten that gap. Nerva–Antonine clears the median by a wide margin; Gordian, despite being the most common dynasty label, trails it.

Use the charts as three separate instruments — name leaders, dynasty spreads, and concentration — rather than collapsing them into a single verdict about who ruled best. Longevity, frequency, and aggregate share each tell a different part of the succession story.

Sources

Data Science Learning Community. (2019). *TidyTuesday: Roman Emperors*. <https://raw.githubusercontent.com/rfordatascience/tidyuesday/main/data/2019/2019-08-13/emperors.csv>

Data Science Learning Community. (2019). *TidyTuesday Roman Emperors source folder*. <https://github.com/rfordatascience/tidyuesday/tree/main/data/2019/2019-08-13>

Boatwright, M. T., Gargola, D. J., Lenski, N., & Talbert, R. J. A. (2012). *The Romans: From Village to Empire* (2nd ed.). Oxford University Press.

Potter, D. S. (2014). *The Roman Empire at Bay, AD 180–395* (2nd ed.). Routledge. <https://doi.org/10.4324/9780203080971>

Ando, C. (2012). *Imperial Rome AD 193 to 284: The Critical Century*. Edinburgh University Press. <https://edinburghuniversitypress.com/book-imperial-rome-ad-193-to-284.html>

Editor's note

Artometrics data report from the TidyTuesday research pipeline. Charts and aggregates are reproducible from the embedded exhibits and public source files.

[View TidyTuesday source on GitHub](#)